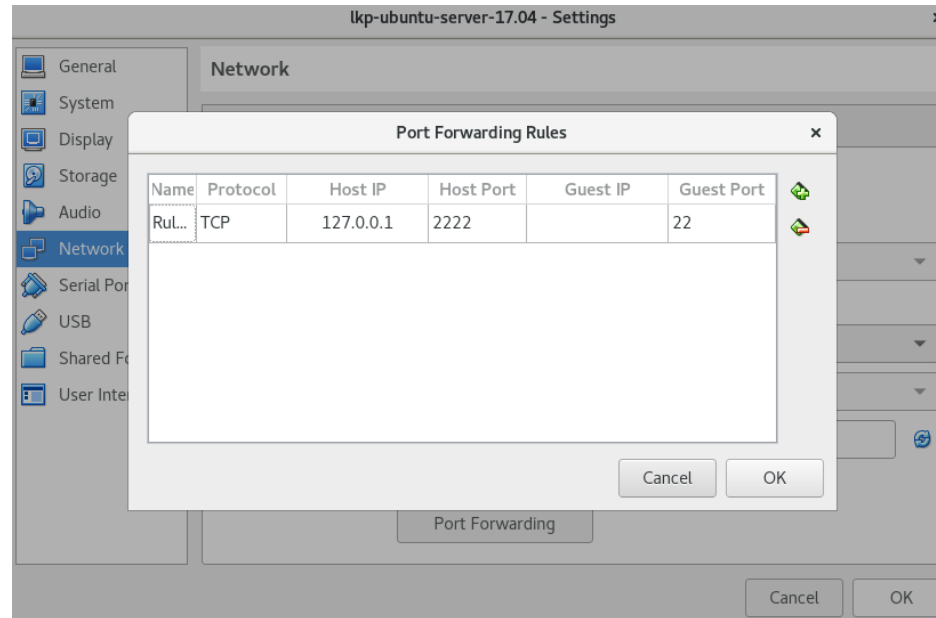


Building and exploring Linux kernel

Dongyoon Lee

Did you succeed in installing Linux on VirtualBox?



- Add a port forwarding rule on VirtualBox

Why software tools are important?

- Linux source code is huge and evolves very fast
 - 27 million lines of code (LoC) ← 1,600 developers / release

```
✓ ~/workspace/research/linux [v5.15]
```

```
$ tree .
```

```
├── arch
│   ├── alpha
│   │   └── boot
│   │       ├── bootloader.lds
│   │       └── bootp.c
│   └── ...
├── lib
│   ├── irqbypass.c
│   ├── Kconfig
│   └── Makefile
└── Makefile
```

7468 directories, 75689 files

Today's lecture

- Tools
 - Version control : `git` , `tig`
 - Configure, build, and install the kernel : `make`
 - Explore the code : `cscope` , `ctags`
 - Editor : `vim` , `emacs`
 - Screen : `tmux`
- Kernel vs. user programming

Obtaining the kernel source code

- Tar ball
 - <https://www.kernel.org/>
- Linus's git repository
 - <git://git.kernel.org/pub/scm/linux/kernel/git/torvalds/linux.git>
- Github mirror of Linus's git repository
 - <https://github.com/torvalds/linux.git>
- Let's explore above web sites!

Version control: `git`

- **Git** is a version control software
 - tracking changes in computer files
- Initially developed by Linus Torvalds for development of the Linux kernel
 - Extensively using in many other software development
 - Github <https://github.com/> is a `git` service provider
- Distributed revision control system
 - Every git directory on every computer is a full-fledged repository with complete history

Essential **git** commands

```
$ # 1. install and configure
$ sudo dnf install git # sudo apt-get install git
$ git config --global user.name "John Doe" # set your name and email for history
$ git config --global user.email johndoe@example.com

$ # 2. creat a repository
$ git init # create a new local repo
$ git clone https://github.com/torvalds/linux.git # clone an existing repo

$ # 3. tags
$ git tag # list all existing tags
$ git checkout v5.15 # checkout the tagged version

$ # 4. commit history (or use tig for prettier output)
$ git log # show all commit history
$ git log <file> # show changes over time for a file
$ git blame <file> # who changed what and when in <file>
```

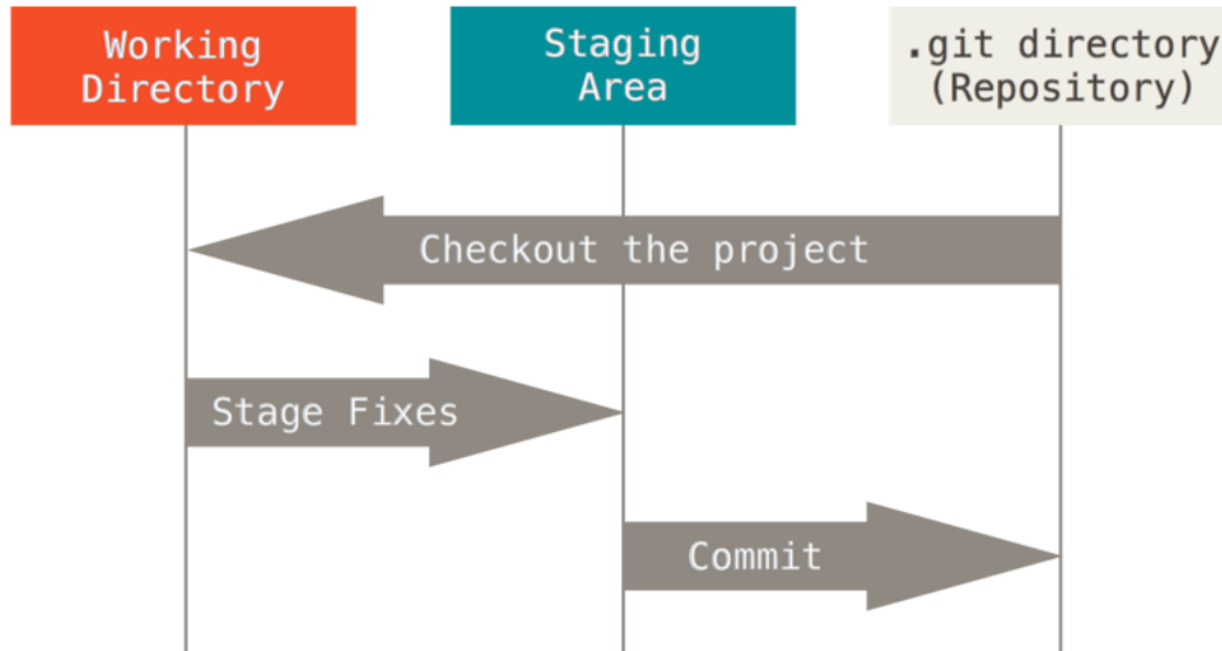
Essential **git** commands

```
$ # 5. local changes
$ git status      # show changed files
$ git diff        # show changed lines
$ git add <file>  # add <file> to the next commit
$ git commit      # commit previously staged files to my local repo

$ # 6. publish and update
$ git push        # publish a committed local changes to a remote repo
$ git pull        # update a local repo
```

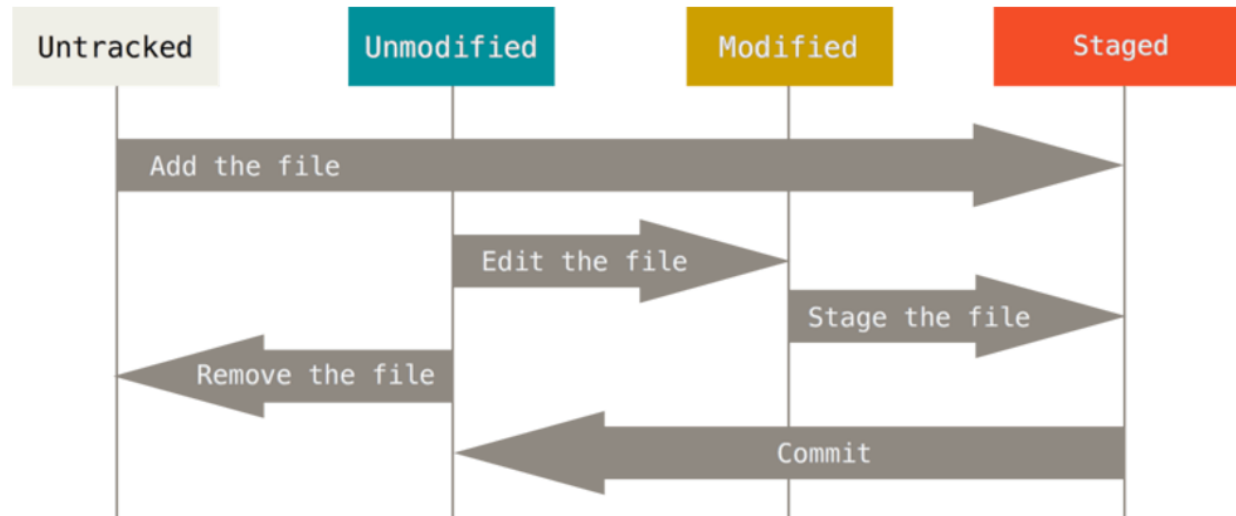
- Many useful git tutorials:
 - [Atlassian](#), [Github](#), [TutorialsPoint](#), [Linux kernel](#), [Pro Git](#)

git workflow



- Source: [Pro Git](#)

git workflow



- Source: [Pro Git](#)

The kernel source tree

```
$ git clone https://github.com/torvalds/linux.git # clone the kernel repo
$ cd linux; git checkout v5.15 # checkout v5.15
$
$ tree -d -L 2      # list top two-level directories
├── arch             # * architecture dependent code
│   ├── arm         #   - ARM architecture
│   └── x86         #   - Intel/AMD x86 architecture
├── block           # * block layer: e.g., IO scheduler
├── Documentation  # * design documents
├── drivers         # * device drivers
│   └── nvme        #   - NVMe SSD
├── fs              # * virtual file system (VFS)
│   ├── ext4       #   - ext4 file system
│   └── xfs         #   - XFS file system
├── include         # * include files
│   ├── linux      #   - include files for kernel
│   └── uapi       #   - include files for user-space tools
├── init            # * bootig: start_kernel() at main.c
├── ipc            # * IPC: e.g., semaphore
└── ...
```

The kernel source tree

```
| ...  
|— kernel      # * core features of the kernel  
|   |— locking  #   - locking: e.g., semaphore, mutex, spinlock  
|   |— sched    #   - task scheduler  
|— lib          # * common library: e.g., red-black tree  
|— mm           # * memory management: e.g., memory allocation, paging  
|— net          # * network stack  
|   |— ipv4      #   - TCP/IPv4  
|   |— ipv6      #   - TCP/IPv6  
|— security     # * security framework  
|   |— selinux   #   - selinux  
|— tools        # * user-space tools  
|   |— perf      #   - perf: performance profiling tool  
|— virt         # * virtualization  
|   |— kvm       #   - KVM type-2 hypervisor
```

615 directories

Build the kernel

1. Configuring the kernel

- Configuration file defining compilation options (~ 3700 for x86)

2. Compiling the kernel

- Compile and link the kernel source code

3. Installing the new kernel

- Install compiled new kernel image to a system
- `make help` to see other make options
- Ref: [Documentation/admin-guide/README.rst](#)

Configure the kernel

- `make menuconfig`
 - Need *libncurses*
 - `sudo dnf install ncurses-devel # Fedora/CentOS/RedHat`
 - `sudo apt-get install libncurses5-dev # Debian/Ubuntu`

Configure the kernel

- `make defconfig`
 - Generate the default configuration of running platform
 - `linux/arch/x86/configs/x86_64_defconfig`
- `make oldconfig`
 - Use the configuration file of running kernel
 - Will ask about new configurations
 - If you are not sure, choose default options
- `make localmodconfig`
 - Update current config disabling modules not loaded

Kernel configuration file: **.config**

- **.config** file is at the root of the kernel source
 - preprocessor flags in the source code

```
# linux/.config
# CONFIG_XEN_PV is not set
CONFIG_KVM_GUEST=y
CONFIG_XFS_FS=m
```

```
/* linux/arch/x86/kernel/cpu/hypervisor.c */
static const __initconst struct hypervisor_x86 * const hypervisors[] =
{
#ifdef CONFIG_XEN_PV
    &x86_hyper_xen_pv,
#endif
#ifdef CONFIG_KVM_GUEST
    &x86_hyper_kvm,
#endif
};
```


Compile the kernel

1. Compile the kernel: `make`

- Compile the kernel source code
- Compiled kernel image: `linux/arch.x86/boot/bzImage`

2. Compile modules: `make modules`

- Parallel `make`
 - `make <target> -j<number of CPUs to use>`
 - E.g., `make -j4`

Install the new kernel

```
# Install the new kernel modules (if you change modules)
```

```
$ sudo make modules_install
```

```
$ ls /lib/modules/
```

```
# Install the new kernel image
```

```
$ sudo make install
```

```
$ ls /boot/*5.15*
```

```
/boot/config-5.15.0      /boot/initrd.img-5.15.0
```

```
/boot/System.map-5.15.0  /boot/vmlinuz-5.15.0
```

```
# Reboot the machine
```

```
$ sudo reboot
```

```
# Check if a system boots with the new kernel
```

```
$ uname -a
```

```
Linux dongyoon 5.15.0 #1 SMP Mon Jan 25 22:56:41 UTC 2021 x86_64 x86_64 x86_64 GNU/Linux
```

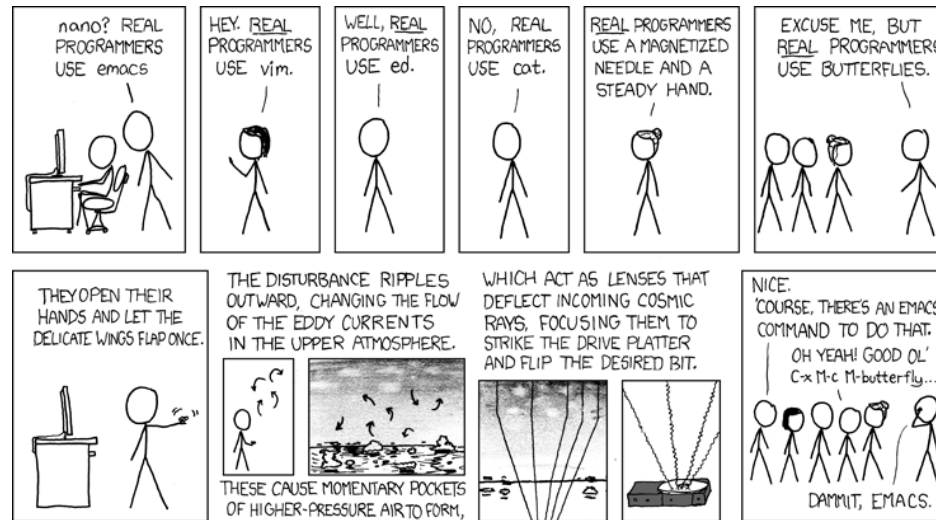
```
# See kernel log
```

```
$ dmesg
```

```
$ dmesg -w # wait for new kernel messages
```

Editor

- There are many good editors
 - vim, emacs



- Source: <https://xkcd.com/378/>


Exploring the code

- [Linux Cross Reference \(LXR\)](#)
- `cscope`
- `vim` with `cscope` or `ctags`
- `emacs` with `cscope`
- ...

Linux Cross Reference (LXR)

- Code indexing tool with a web interface
 - Don't install it! One instance is running here:
 - <http://lxr.free-electrons.com/>
- Allows to:
 - Browse the code of different Linux versions
 - Search for identifiers (functions, variables, etc.)
 - Quickly lookup a function declaration/definition

Linux Cross Reference (LXR)



Linux Cross Reference

Free Electrons
Embedded Linux Experts

[Source Navigation](#) • [Identifier Search](#) • [Freetext Search](#) •

Version: 2.0.40 2.2.26 2.4.37 3.12 3.13 3.14 3.15 3.16 3.17 3.18 3.19 **4.0** 4.1 4.2 4.3 4.4 4.5 4.6 4.7 4.8 4.9

Linux/fs/

Parent directory

9p/

adfs/

affs/

afs/

autofs4/

befs/

bfs/

btrfs/

cachefiles/

ceph/

cifs/

coda/

configs/

cramfs/

debuafs/

Linux Cross Reference (LXR)

Linux Cross Reference (LXR)

```

533         static_command_line, __start_param,
534         __stop_param - __start_param,
535         -1, -1, &unknown_bootoption);
536 if (!IS_ERR_OR_NULL(after_dashes))
537     parse_args("Setting init args", after_dashes, NULL, 0, -1, -1,
538         set_init_arg);
539
540 jump_label_init();
541
542 /*
543  * These use large bootmem allocations and must precede
544  * kmem_cache_init()
545  */
546 setup_log_buf(0);
547 pidhash_init();
548 vfs_caches_init_early();
549 sort_main_extable();
550 trap_init();
551 mm_init();
552
553 /*
554  * Set up the scheduler prior starting any interrupts (such as the
555  * timer interrupt). Full topology setup happens at smp_init()
556  * time - but meanwhile we still have a functioning scheduler.
557  */
558 sched_init();
559 /*
560  * Disable preemption - early bootup scheduling is extremely
561  * fragile until we cpu_idle() for the first time.
562  */
563 preempt_disable();
564 if (WARN(!irqs_disabled(),
565     "Interrupts were enabled *very* early, fixing it\n"))
566     local_irq_disable();
567 idr_init_cache();
568 rcu_init();
569
570 /* trace_printk() and trace points may be used after this */
571 trace_init();
572
573 context_tracking_init();
574 radix_tree_init();
575 /* init some links before init_ISA_irqs() */

```

Click on a
function call to
search
for the function

Linux Cross Reference (LXR)

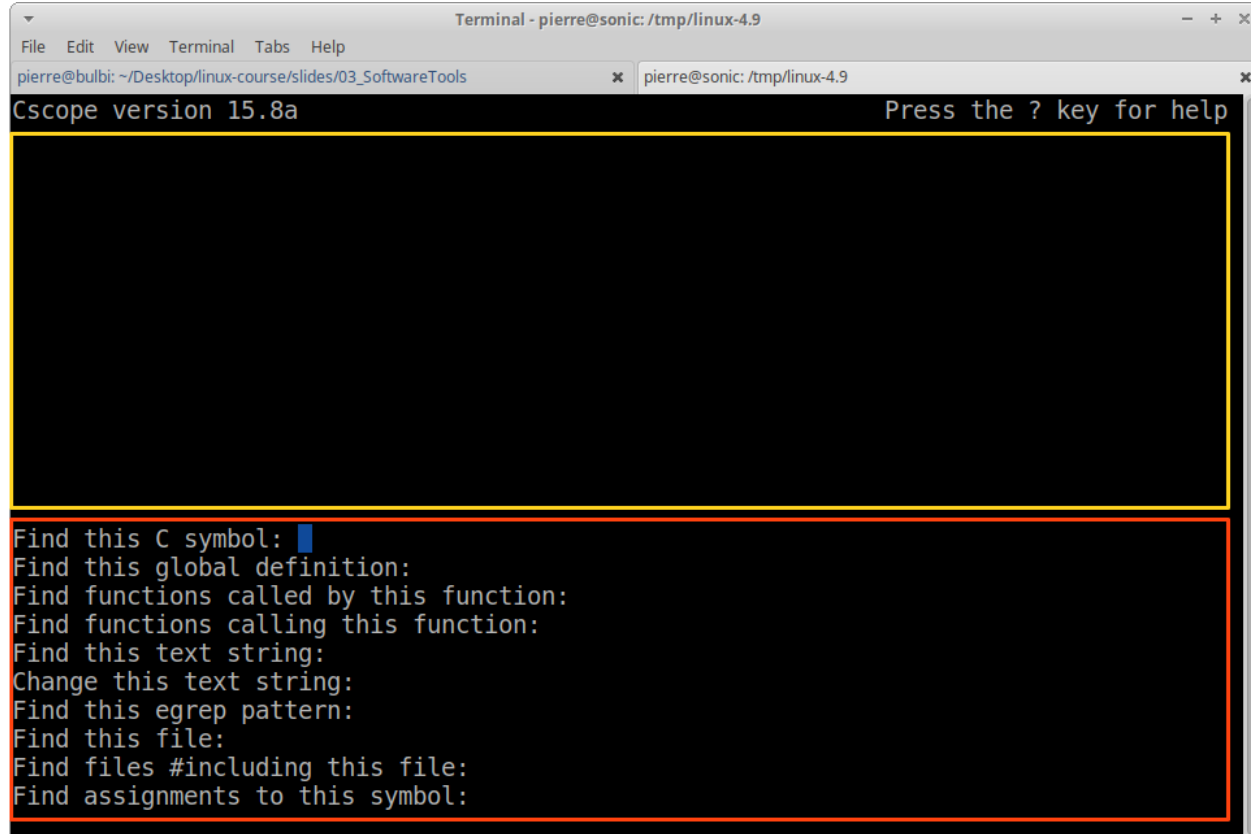
cscope

- Command line tool to browse (potentially large) C codebases
- Installation: `sudo {apt-get|dnf} install cscope`
- Build cscope database
 - `cd linux; KBUILD_ABS_SRCTREE=1 ARCH=x86 make cscope` # only for x86
 - Need to rebuild after code changes
- Although `cscope -R` is the common way to build cscope database, `make cscope` is optimized for the kernel source code

cscope

- Search for:
 - C identifier occurrences (variable name, function name, typedef/struct, label)
 - Functions/variables definitions
 - Functions called by/calling function f
 - Text string
- Terminating cscope: Ctrl-d

cscope



```
Terminal - pierre@sonic: /tmp/linux-4.9
File Edit View Terminal Tabs Help
pierre@bulbi: ~/Desktop/linux-course/slides/03_SoftwareTools x pierre@sonic: /tmp/linux-4.9 x
Cscope version 15.8a Press the ? key for help

Find this C symbol: 
Find this global definition:
Find functions called by this function:
Find functions calling this function:
Find this text string:
Change this text string:
Find this egrep pattern:
Find this file:
Find files #including this file:
Find assignments to this symbol:
```

cscope

```
Terminal - pierre@sonic: /tmp/linux-4.9
File Edit View Terminal Tabs Help
pierre@bulbi: ~/Desktop/linux-course/slides/03_SoftwareTools x pierre@sonic: /tmp/linux-4.9 x
Cscope version 15.8a Press the ? key for help

Find this C symbol: spin lock
Find this global definition:
Find functions called by this function:
Find functions calling this function:
Find this text string:
Change this text string:
Find this egrep pattern:
Find this file:
Find files #including this file:
Find assignments to this symbol:
```

cscope

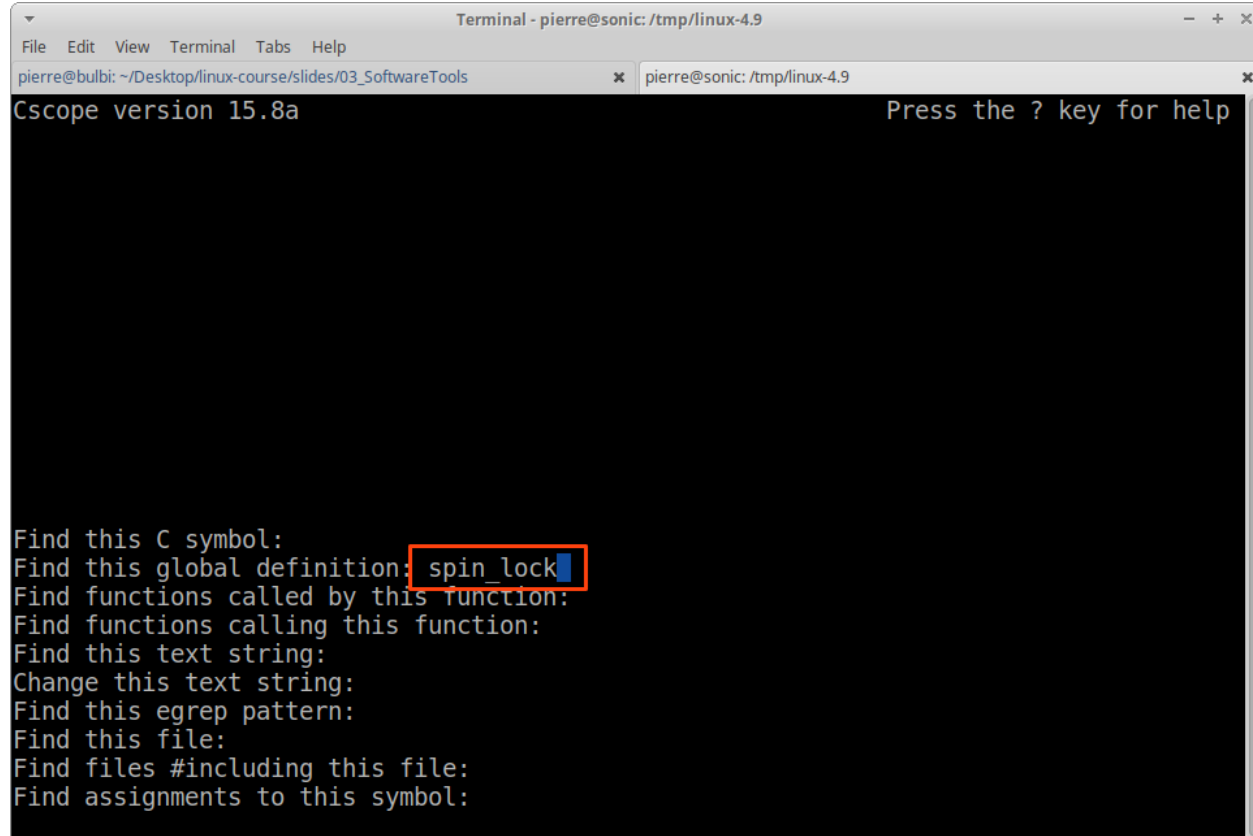
```
Terminal - pierre@sonic: /tmp/linux-4.9
File Edit View Terminal Tabs Help
pierre@bulbi: ~/Desktop/linux-course/slides/03_SoftwareTools x pierre@sonic: /tmp/linux-4.9 x
C symbol: spin_lock

File           Function      Line
0 platsmp.c     <global>      287 spin_lock(&boot_lock);
1 bus.c         <global>      1049 spin_lock(&device_klist->k_lock);
2 platform.c    <global>      697 spin_lock(&drv->driver->.bus->p->klist_drivers->.k_lock);
3 runtime.c     <global>      279 spin_lock(&dev->power->.lock);
4 omap_gem.c    <global>      1256 spin_lock(&sync_lock);

* Lines 1-6 of 9743, 9738 more - press the space bar to display more *

Find this C symbol:
Find this global definition:
Find functions called by this function:
Find functions calling this function:
Find this text string:
Change this text string:
Find this egrep pattern:
Find this file:
Find files #including this file:
Find assignments to this symbol:
```

cscope



```
Terminal - pierre@sonic: /tmp/linux-4.9
File Edit View Terminal Tabs Help
pierre@bulbi: ~/Desktop/linux-course/slides/03_SoftwareTools x pierre@sonic: /tmp/linux-4.9 x
Cscope version 15.8a Press the ? key for help

Find this C symbol:
Find this global definition: spin_lock
Find functions called by this function:
Find functions calling this function:
Find this text string:
Change this text string:
Find this egrep pattern:
Find this file:
Find files #including this file:
Find assignments to this symbol:
```

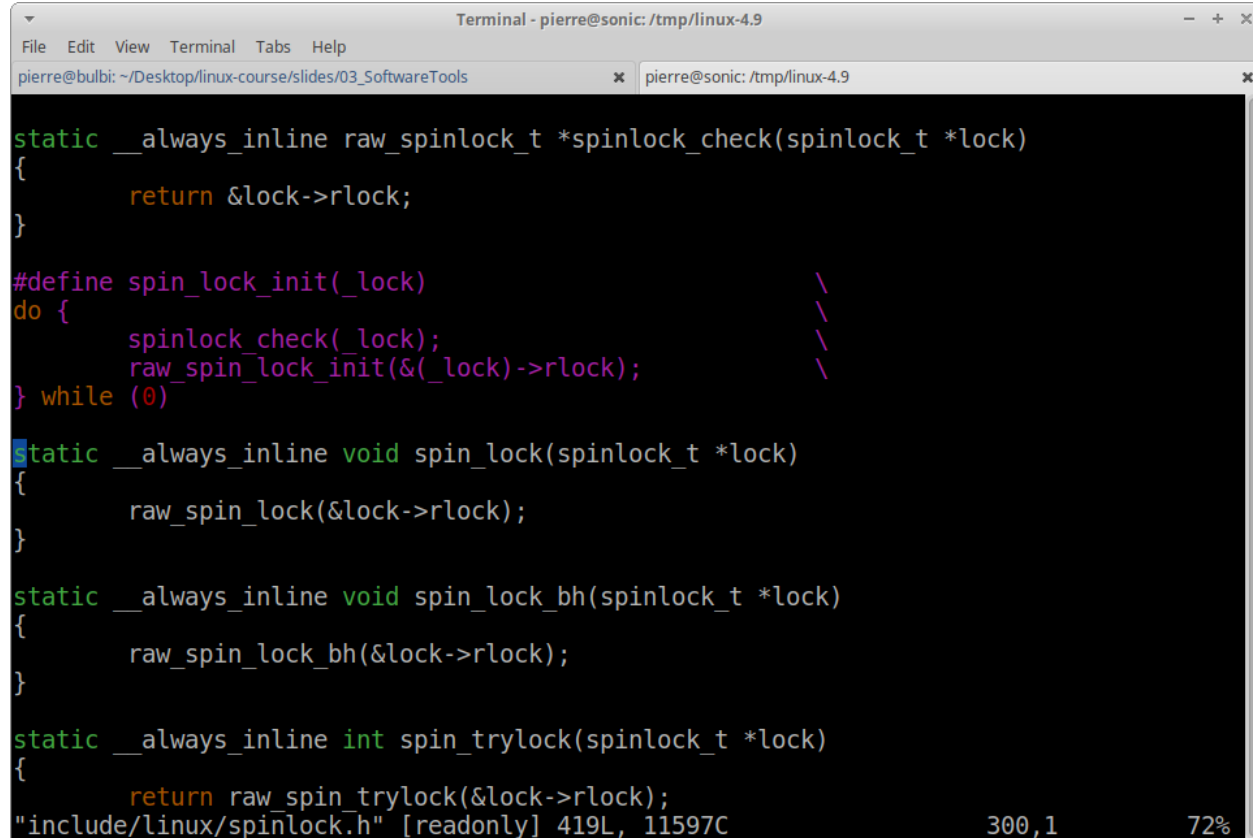
cscope

```
Terminal - pierre@sonic: /tmp/linux-4.9
File Edit View Terminal Tabs Help
pierre@bulbi: ~/Desktop/linux-course/slides/03_SoftwareTools x pierre@sonic: /tmp/linux-4.9 x
Global definition: spin_lock

File      Line
0 aic79xx_osm.h 352 spinlock_t spin_lock;
1 aic7xxx_osm.h 356 spinlock_t spin_lock;
2 comedidev.h  177 spinlock_t spin_lock;
3 spinlock.h    300 static always inline void spin_lock(spinlock_t *lock)

Find this C symbol:
Find this global definition:
Find functions called by this function:
Find functions calling this function:
Find this text string:
Change this text string:
Find this egrep pattern:
Find this file:
Find files #including this file:
Find assignments to this symbol:
```


cscope

A terminal window titled "Terminal - pierre@sonic: /tmp/linux-4.9" displays C code for spinlocks. The code includes functions like raw_spinlock_check, spin_lock_init, spin_lock, spin_lock_bh, and spin_trylock. The terminal window has a menu bar with "File", "Edit", "View", "Terminal", "Tabs", and "Help". The status bar at the bottom shows "300,1" and "72%".

```
static __always_inline raw_spinlock_t *spinlock_check(spinlock_t *lock)
{
    return &lock->rlock;
}

#define spin_lock_init(_lock) \
do { \
    spinlock_check( _lock); \
    raw_spin_lock_init(&(_lock)->rlock); \
} while (0)

static __always_inline void spin_lock(spinlock_t *lock)
{
    raw_spin_lock(&lock->rlock);
}

static __always_inline void spin_lock_bh(spinlock_t *lock)
{
    raw_spin_lock_bh(&lock->rlock);
}

static __always_inline int spin_trylock(spinlock_t *lock)
{
    return raw_spin_trylock(&lock->rlock);
}

#include/linux/spinlock.h" [readonly] 419L, 11597C
```

vim with cscope or ctags

- vim can use the tag database of cscope ,as well as ctags
 - `sudo apt-get install cscope exuberant-ctags`
 - `sudo dnf install cscope ctags`
- Generate the database
 - `cd linux; make cscope tags -j2`
- Launch vim
 - `vim init/main.c`

vim with cscope or ctags

- Search for function definition/variable declaration:
 - `:tag start_kernel` or `:cs find global start_kernel`
- Help for `ctags` and `cscope`
 - `:help tag` or `:help cs`
- Another way to find a function definition/variable declaration:
 - Put the cursor on the symbol name and press `Ctrl+]`
- To navigate back and forth between file:
 - `:bp` or `:bn`

Screen: tmux

- tmux is a tool to manage virtual consoles

The image displays two terminal windows side-by-side, both titled 'changwoo:changwoo:changwoo'.

The left terminal window shows the compilation and execution of a C program. The code includes standard Linux headers and a main function. The output shows the program running successfully. At the bottom, a system status bar is visible, showing the current directory as '/workspace/research/linux' and the time as 'Sun 20 Aug 14:22'.

The right terminal window shows a search for the configuration option 'CONFIG_KVM_GUEST' in the kernel headers. The search is performed using the 'grep' command. The output shows the configuration option is defined in the 'autoconf.h' file, indicating that KVM guest support is enabled.

- Ref: [A tmux Primer](#)

Essential **tmux** commands

- **tmux** : start a new tmux session
- **Ctrl-b %** : split a pane vertically
- **Ctrl-b "** : split a pane horizontally
- **Ctrl-b o** : move to the next pane
- **Ctrl-b z** : zoom (or unzoom) a pane
- **Ctrl-b c** : create a new window
- **Ctrl-b N** : go to window N (0~9)
- **Ctrl-b d** : detach from a session
- **tmux a** : attach to an existing session

Kernel vs. user programming

- No libc or standard headers
 - Instead the kernel implements lots of libc-like functions
- Examples
 - `#include <string.h>` → `#include <linux/string.h>`
 - `printf("Hello!")` → `printk(KERN_INFO "Hello!")`
 - `malloc(64)` → `kmalloc(64, GFP_KERNEL)`

Kernel vs. user programming

- Use GCC extensions
- Inline functions
 - `static inline void func()`
- Inline assembly: less than 2%
 - `asm volatile("rdtsc" : "=a" (l), "=d" (h));`
- Branch annotation: hint for better optimization
 - `if (unlikely(error)) {...}`
 - `if (likely(success)) {...}`

Kernel vs. user programming

- No (easy) use of floating point
- Small, fixed-size stack: 8 KB (2 pages) in x86
- No memory protection
 - SIGSEGV → **kernel panic (oops)**
- [An example of kernel oops](#)

Kernel vs. user programming

- Synchronization and concurrency
 - Multi-core processor → synchronization among tasks
 - A kernel code can execute on two more processors
 - Preemptive multitasking → synchronization among tasks
 - A task can be scheduled and re-scheduled at any time
 - Interrupt → synchronization with interrupt handlers
 - Can occur in the midst of execution (e.g., accessing resource)
 - Need to synchronize with interrupt handler

Linux kernel coding style

- Indentation: 1 tab → 8-character width (not 8 spaces)
- No CamelCase use underscores: `SpinLock` → `spin_lock`
- Use C-style comments: `/* use this style */` not this
- Line length: 80 column
- **Write code in a similar style with other kernel code**
- Ref: <Documentation/process/coding-style.rst>

Linux kernel coding style

```
/*  
 * a multi-lines comment  
 * (no C++ '//' !)  
 */  
  
struct foo {  
    int member1;  
    double member2;  
}; /* no typedef ! */  
  
#ifdef CONFIG_COOL_OPTION  
int cool_function(void) {  
    return 42;  
}  
#else  
int cool_function(void) { }  
#endif /* CONFIG_COOL_OPTION */
```

Linux kernel coding style

```
void my_function(int the_param, char *string, int a_long_parameter,
                 int another_long_parameter)
{
    int x = the_param % 42;

    if (!the_param)
        do_stuff();

    switch (x % 3) {
    case 0:
        do_some_stuff();
        cool_function();
        break;
    case 1:
        /* Fall through */
    default:
        do_other_stuff();
        cool_function();
    }
}
```

Summary of tools

- Version control : `git`, `tig`
- Configure the kernel : `make oldconfig`
- Build the kernel : `make -j8; make modules -j8`
- Install the kernel : `make install; make modules_install`
- Explore the code : `make cscope tags -j2; cscope, ctags`
- Editor : `vim`, `emacs`
- Screen : `tmux`

Other useful sources II

- [Documentation directory](#): the most up-to-date design documents
- [The Linux Kernel Documentation](#): the extensive documents extracted from kernel source
- [Linux Weekly News](#): easy explanation of recently added kernel features
- [Linux Inside](#): textbook-style description on kernel subsystems
- [Kernel newbies](#): useful information for new kernel developers
- [Linux Kernel API Manual](#)
- [Kernel Recipes](#)
- [kernel planet](#)

Next actions

- Master the essential tools, seriously
 - editor: `vim`, `emacs`
 - code navigation: `cscope`, `ctags`
 - version control: `git`, `tig`
 - terminal: `ssh`, `tmux`
- Useful lecture videos: [Vim](#), [tmux](#), [ssh](#), [Git](#)

Next lecture

- Isolation and system call
- Explore how following three system calls are implemented in the kernel

```
fd = open("out", 1);  
write(fd, "hello\n", 6);  
pid = fork();
```