

Kernel debugging

Dongyoon Lee

Summary of last lectures

- Kernel data structures
 - list, hash table, red-black tree
 - radix tree, XArray, bitmap array
- Kernel module

Today's agenda

- Kernel debugging
 - tools, techniques, and tricks

Kernel development cycle

- Write code → Build kernel/modules → Deploy → **Test and debug**
- *Debugging is the real bottleneck* even for experienced kernel developers due to limitations in kernel debugging
- It is important to get used to kernel debugging techniques to save your time and effort

Kernel debugging techniques

- Print debug message: `printk`
- Assert your code: `BUG_ON(c)`, `WARN_ON(c)`
- Analyze kernel panic message
- Debug with QEMU/gdb

Print debug message: **printk**

- Similar to `printf()` in C library
- Need to specify a log level (the default level is `KERN_WARNING` or `KERN_ERR`)

```
KERN_EMERG    /* 0: system is unusable          */
KERN_ALERT    /* 1: action must be taken immediately        */
KERN_CRIT     /* 2: critical conditions                     */
KERN_ERR      /* 3: error conditions                       */
KERN_WARNING  /* 4: warning conditions                     */
KERN_NOTICE   /* 5: normal but significant condition       */
KERN_INFO     /* 6: informational                         */
KERN_DEBUG    /* 7: debug-level messages                   */
```

e.g., `printk(KERN_DEBUG "debug message from %s:%d\n", __func__, __LINE__);`

Print debug message: **printk**

- Prints out only messages, which log level is higher than the current level.

```
# Check current kernel log level
$ cat /proc/sys/kernel/printk
    4      4      1      7
#  |      |      |      |
# current default minimum boot-time-default
# Enable all levels of messages:
$ echo 7 > /proc/sys/kernel/printk
```

- The kernel message buffer is a fixed-size circular buffer.
- If the buffer fills up, it warps around and *you can lose some message*.
- Increasing the buffer size would be helpful a little bit.
 - Add `log_buf_len=1M` to kernel boot parameters (power of 2)

Print debug message: **printk**

- Support additional format specifiers

```
/* function pointers with function name */
"%pF"    versatile_init+0x0/0x110 /* symbol+offset/length */
"%pf"    versatile_init

/* direct code address (e.g., regs->ip) */
"%pS"    versatile_init+0x0/0x110
"%ps"    versatile_init

/* direct code address in stack (e.g., return address) */
"%pB"    prev_fn_of_versatile_init+0x88/0x88

/* Example */
printk("Going to call: %pF\n", p->func);
printk("Faulted at %pS\n", (void *)regs->ip);
printk(" %s%pB\n", (reliable ? "" : "? "), (void *)*stack);
```

- Ref: [How to get printk format specifiers right](#)

BUG_ON(c), WARN_ON(c)

- Similar to `assert(c)` in userspace
- `BUG_ON(c)`
 - if `c` is false, kernel panics with its call stack
- `WARN_ON(c)`
 - if `c` is false, kernel prints out its call stack and keeps running

Kernel panic message

```
[ 174.507084] Stack:
[ 174.507163] ce0bd8ac 00000008 00000000 ce4a7e90 c039ce30 ce0bd8ac c0718b04 c07185a0
[ 174.507380] ce4a7ea0 c0398f22 ce0bd8ac c0718b04 ce4a7eb0 c037deee ce0bd8e0 ce0bd8ac
[ 174.507597] ce4a7ec0 c037dfe0 c07185a0 ce0bd8ac ce4a7ed4 c037d353 ce0bd8ac ce0bd8ac
[ 174.507888] Call Trace:
[ 174.508125] [] ? sd_remove+0x20/0x70
[ 174.508235] [] ? scsi_bus_remove+0x32/0x40
[ 174.508326] [] ? __device_release_driver+0x3e/0x70
[ 174.508421] [] ? device_release_driver+0x20/0x40
[ 174.508514] [] ? bus_remove_device+0x73/0x90
[ 174.508606] [] ? device_del+0xef/0x150
[ 174.508693] [] ? __scsi_remove_device+0x47/0x80
[ 174.508786] [] ? scsi_remove_device+0x22/0x40
[ 174.508877] [] ? __scsi_remove_target+0x94/0xd0
[ 174.508969] [] ? __remove_child+0x0/0x20
[ 174.509060] [] ? __remove_child+0x17/0x20
[ 174.509148] [] ? device_for_each_child+0x38/0x60
```

- **Q: How can I find where `sd_remove+0x20/0x70` is at source code?**

Analyze kernel panic message

1. Find where `sd_remove()` is (e.g., `linux/driver/scsi/sd.c`)
2. Load its object file with `gdb`
3. Use `gdb` command, `list *(function+offset)` command

```
manjo@hungry:~/devel/kernel/build/build-generic/drivers/scsi$ gdb sd.o
This GDB was configured as "x86_64-linux-gnu".
(gdb) list *(sd_remove+0x20)
0x1650 is in sd_remove (/home/manjo/devel/linux/drivers/scsi/sd.c:2125).
2120     static int sd_remove(struct device *dev)
2121     {
2122         struct scsi_disk *sdkp;
2123
2124         async_synchronize_full();
```

- **Q: Can I debug kernel using `gdb`? → It is possible using `QEMU/gdb`**

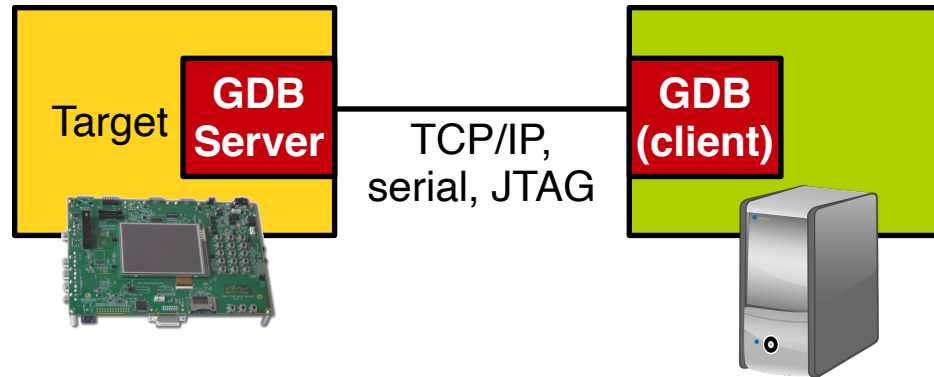
QEMU

- Full system emulator: emulates an entire virtual machine
 - Using a software model for the CPU, memory, devices
 - Emulation is slow
- Can also be used in conjunction with hardware virtualization extensions to provide high performance virtualization
 - **KVM**: In-kernel support for virtualization + extensions to QEMU
- To install QEMU, libvirt, GDB, etc.


```
$ sudo apt install qemu qemu-system qemu-kvm libvirt-daemon-system \
    libvirt-clients bridge-utils gdb
```

GDB server

- Originally used to debug a program executing on a remote machine
- For example when GDB is not available on that remote machine
 - E.g., low performance embedded systems



Kernel debugging with QEMU/gdb



- Linux kernel runs in a virtual machine (KVM or emulated on QEMU)
- Hardware devices are emulated with QEMU
- GDB server runs at QEMU, emulated virtual machine, so it can fully control Linux kernel running on QEMU
- *It is fantastic for debugging and code exploration!*

Three steps for QEMU/gdb

1. Build kernel for QEMU/gdb debugging
2. Create the root filesystem
3. Run Linux on QEMU and attach GDB

Build kernel for QEMU/gdb debugging

- Rebuild kernel with debug info, gdb script, and virtio enabled
- Following should be built-in, not built as a kernel module.

```
$ scripts/config -e CONFIG_DEBUG_INFO
$ scripts/config -e CONFIG_GDB_SCRIPTS
$ scripts/config -e CONFIG_E1000
$ scripts/config -e CONFIG_VIRTIO
$ scripts/config -e CONFIG_VIRTIO_NET
$ scripts/config -e CONFIG_VIRTIO_BLK

$ cat .config
CONFIG_DEBUG_INFO=y          # debug symbol
CONFIG_GDB_SCRIPTS=y         # qemu/gdb support
CONFIG_E1000=y               # default network card

CONFIG_VIRTIO=y              # virtio
CONFIG_VIRTIO_NET=y          # virtio network
CONFIG_VIRTIO_BLK=y          # virtio block device
```


Build kernel for QEMU/gdb debugging

```
# or (use `/ GDB_SCRIPTS` in `make menuconfig`  
$ make menuconfig  
| Prompt: Provide GDB scripts for kernel debugging  
| Location:  
|   -> Kernel hacking  
|     -> Compile-time checks and compiler options  
|       -> Provide GDB scripts for kernel debugging  
|         Compile the kernel with frame pointers  
  
# then build the kernel  
$ make -j8; make -j8 modules
```

- You don't need `make modules_install; make install` because all necessary features will be built-in to ease of deployment.

Build kernel for QEMU/gdb debugging

- You can find the following files under the Linux source folder:

```
[path-to-linux]/arch/x86/boot/bzImage    # kernel binary image  
[path-to-linux]/vmlinux                  # target for GDB  
[path-to-linux]/vmlinux-gdb.py           # GDB helper script
```

- Add *vmlinux-gdb.py* script to GDB's auto load path so that you can later use the lx- helper commands.

```
$ echo "add-auto-load-safe-path [path-to-linux]/vmlinux-gdb.py" >> ~/.gdbinit
```

Create the root filesystem

- We need a root filesystem for the kernel to boot on QEMU.
- Two options
 - Create a new image using the [buildroot](#) project.
 - Convert a VirtualBox image (.vdi) to a QEMU (.img). See [this](#).

Create the root filesystem with buildroot

Clone the buildroot project

```
$ git clone git://git.buildroot.net/buildroot
```

Configure buildroot to include the following options

```
$ cd buildroot
```

```
$ make menuconfig
```

```
Target options -- Target architecture -- select [x86_64]
```

```
Toolchain -- Enable C++ support [*]
```

```
Filesystem images -- ext2/3/4 root filesystem -- choose [ext4]
```

```
Target packages -- Network applications -- select openssh [*]
```

Build

```
$ make -j<number of CPUs>
```

Find the output image below

```
$ ls output/images/rootfs.ext4
```

```
output/images/rootfs.ext4
```

Run kernel with QEMU/gdb

run QEMU/gdb

```
sudo qemu-system-x86_64 \  
-s \                               # enable qemu-gdb debugging  
-S \                               # pause on the first instr.  
-kernel <path-to-linux>/arch/x86_64/boot/bzImage \ # kernel image  
-nographic \  
-drive format=raw,file=<path-to-buildroot>/output/images/rootfs.ext4,if=virtio \  
-append "root=/dev/vda console=ttyS0 nokaslr other-paras-here-if-needed" \  
-m <mem 4G> \  
-cpu host \  
-smp <num of cpus> \  
-net nic,model=virtio \  
-net user,hostfwd=tcp::10022-:22 \ # port forwarding for ssh  
-enable-kvm                       # use kvm
```

Ctrl-a x: terminating QEMU

Run kernel with QEMU/gdb

- QEMU options
 - `-kernel vmlinux` : path to the vmlinux of the kernel to debug
 - `-s` : enable the GDB server and open a port 1234
 - `-S` : (optional) pause on the first kernel instruction waiting for a GDB client connection to continue

Connect to the kernel on QEMU/gdb

- Run gdb and attach to QEMU

```
$ gdb <path-to-linux>/vmlinux  
(gdb) target remote :1234
```

- You can use all *gdb commands* and *Linux-provided gdb helpers!*

- [b]reak <function name or filename:line# or *memory address>
- [h]break <start_kernel or any function name> # to debug boot code
- [d]elete <breakpoint #>
- [c]ontinue
- [b]ack[t]race
- [i]nfo [b]reak
- [n]ext
- [s]tep
- [p]rint <variable or *memory address>
- Ctrl-x Ctrl-a: TUI mode

Linux-provided gdb helpers

- Load module and main kernel symbols

```
(gdb) lx-symbols
loading vmlinux
scanning for modules in /home/user/linux/build
loading @0xfffffffffa0020000: /home/user/linux/build/net/netfilter/xt_tcpudp.ko
loading @0xfffffffffa0016000: /home/user/linux/build/net/netfilter/xt_pkttype.ko
loading @0xfffffffffa0002000: /home/user/linux/build/net/netfilter/xt_limit.ko
loading @0xfffffffffa00ca000: /home/user/linux/build/net/packet/af_packet.ko
loading @0xfffffffffa003c000: /home/user/linux/build/fs/fuse/fuse.ko
...
loading @0xfffffffffa0000000:
/home/user/linux/build/drivers/ata/ata_generic.ko
```


Linux-provided gdb helpers

- Set a breakpoint on some not yet loaded module function, e.g.:

```
(gdb) b btrfs_init_sysfs
Function "btrfs_init_sysfs" not defined.
Make breakpoint pending on future shared library load? (y or [n]) y
Breakpoint 1 (btrfs_init_sysfs) pending.
```

- Continue the target:

```
(gdb) c

loading @0xfffffffffa006e000: /home/user/linux/build/lib/zlib_deflate/zlib_deflate.ko
loading @0xfffffffffa01b1000: /home/user/linux/build/fs/btrfs/btrfs.ko

Breakpoint 1, btrfs_init_sysfs () at /home/user/linux/fs/btrfs/sysfs.c:36
36          btrfs_kset = kset_create_and_add("btrfs", NULL, fs_kobj);
```

Linux-provided gdb helpers

- Dump the log buffer of the target kernel:

```
(gdb) lx-dmesg
[ 0.000000] Initializing cgroup subsys cpuset
[ 0.000000] Initializing cgroup subsys cpu
[ 0.000000] Linux version 3.8.0-rc4-dbg+ (...
[ 0.000000] Command line: root=/dev/sda2 resume=/dev/sda1 vga=0x314
[ 0.000000] e820: BIOS-provided physical RAM map:
[ 0.000000] BIOS-e820: [mem 0x0000000000000000-0x000000000009fbff] usable
[ 0.000000] BIOS-e820: [mem 0x000000000009fc00-0x000000000009ffff] reserved
```

- Examine fields of the current task struct:

```
(gdb) p $lx_current().pid
$1 = 4998
```

Linux-provided gdb helpers

- Help

```
(gdb) apropos lx
function lx_current -- Return current task
function lx_module -- Find module by name and return the module variable
function lx_per_cpu -- Return per-cpu variable
function lx_task_by_pid -- Find Linux task by PID and return the task_struct variable
function lx_thread_info -- Calculate Linux thread_info from task variable
lx-dmesg -- Print Linux kernel log buffer
lx-lsmod -- List currently loaded modules
lx-symbols -- (Re-)load symbols of Linux kernel and currently loaded
modules
```

Tips for QEMU-gdb kernel debugging

- `(gdb) p my_var` → `$1 = <value optimized out>`
 - `my_var` is optimized out
 - Since it is not possible to disable optimization for the entire kernel, we need to disable optimization for a specific file.

```
# linux/fs/ext4/Makefile
obj-$(CONFIG_EXT4_FS) += ext4.o

CFLAGS_bitmap.o = -O0 # disable optimization of bitmap.c

ext4-y := balloc.o bitmap.o dir.o file.o \
#...
```

Tips for QEMU-gdb kernel debugging

- Cursor disappears in qemu window?
 - Ctrl Alt (right)
- Always terminates QEMU with the `halt` command otherwise the disk image could be corrupted
- QEMU is too slow
 - Try KVM (`enable-kvm`)
 - It works only when your host is Linux.

Further readings

- [Debugging by printing](#)
- [Kernel Debugging Tricks](#)
- [Kernel Debugging Tips](#)
- [Debugging kernel and modules via gdb](#)
- [gdb Cheatsheet](#)
- [Speed up your kernel development cycle with QEMU](#)
- [Installing new Debian systems with debootstrap](#)
- [Migrate a VirtualBox Disk Image \(.vdi\) to a QEMU Image \(.img\)](#)
- [The kernel's command-line parameters](#)

Next lecture

- Process management