

OS Basic: Process and Thread

Dongyoon Lee

Summary of last lectures

- Getting, building, and exploring the Linux kernel
- System call: interface between applications and kernel
- Kernel data structures
- Kernel modules
- Kernel debugging techniques

Today's agenda

- Process
- Thread

Process

- A process is a program in execution .
- A process is not the same thing as a program:
 - A program is a *passive* entity.
 - Processes are *active*.
 - Each process only runs one program at a time.
 - The same program can be run by more than one process at a time.

Multiprogramming

- A multiprogramming OS supports many **concurrent** processes.
 - Each process has a **context**, including an *address space*, and can receive *CPU cycles*.
 - The OS achieves an illusion of concurrency by switching the CPU rapidly between processes.

Process Context

- The context of a process is essentially a snapshot of the state of that process, including:
 - The **CPU state**, including contents of CPU registers.
 - The **run state** of the process (running, waiting, ready).
 - The **address space** of the process, which is its “view of memory.”

This includes:

- *Main memory* allocated to the process.
- *Page tables* that describe a mapping from virtual to physical addresses (more later).

(1) CPU state (x86_64)

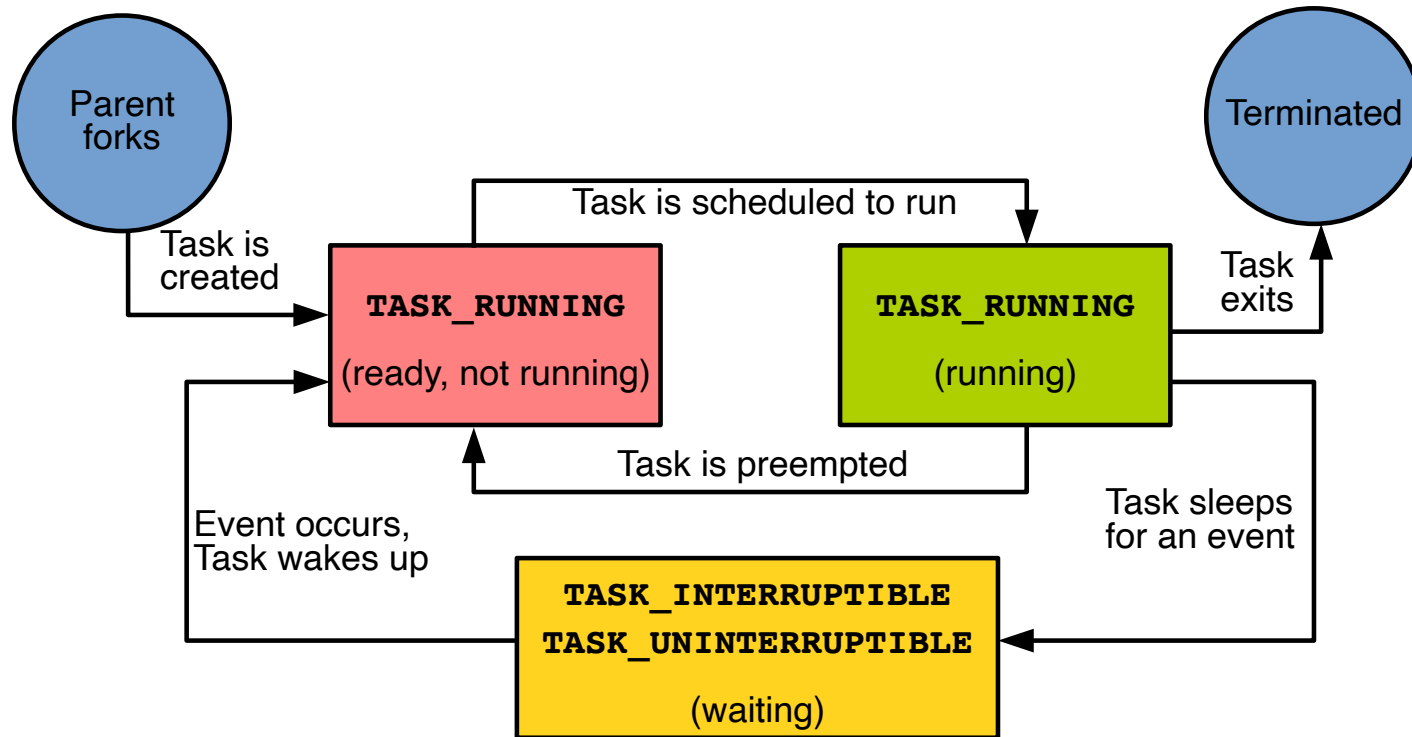
| Register | Purpose | Saved across calls |
|----------|----------------------------------------|--------------------|
| %rax | temp register; return value | No |
| %rbx | callee-saved | Yes |
| %rcx | used to pass 4th argument to functions | No |
| %rdx | used to pass 3rd argument to functions | No |
| %rsi | used to pass 2nd argument to functions | No |
| %rdi | used to pass 1st argument to functions | No |
| %r8 | used to pass 5th argument to functions | No |
| %r9 | used to pass 6th argument to functions | No |

(1) CPU state (x86_64)

| Register | Purpose | Saved across calls |
|----------|----------------------------|--------------------|
| %r10-r11 | temporary | No |
| %r12-r15 | callee-saved registers | Yes |
| %rsp | stack pointer | Yes |
| %rbp | callee-saved; base pointer | Yes |


- Plus, floating-point and SIMD registers
- The program counter (PC) register points to the address of the next instruction to be executed from memory

(2) Run State of a Process



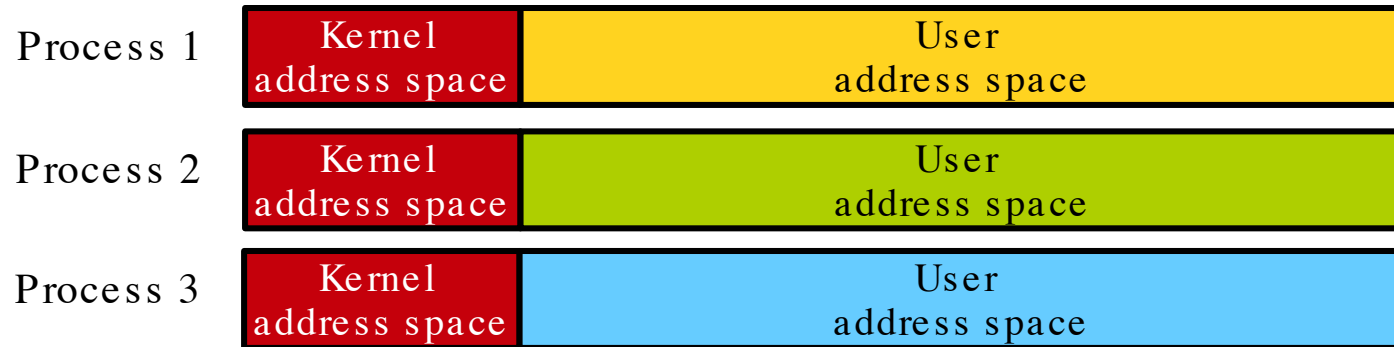
(3) Address Space

- An address space is the “view of memory” provided by the operating system for a process.



Multiple Address Spaces and OS Kernel

- A multiprogrammed OS maintains multiple address spaces simultaneously



Virtual Memory

- Virtual memory is a mechanism that permits a process to be run without having the entire contents of its address space loaded into main memory at one time.
 - The part of the address space that is loaded into main memory is called **resident**.
 - The remainder is called **nonresident**.
- Nonresident data is saved on a secondary storage area, called the **backing store**.

What is Virtual Memory Good For?

- A primary purpose of virtual memory is to increase the degree of multiprogramming to obtain more efficient utilization of system resources:
 - More runnable processes can be kept in main memory at one time.
 - More runnable processes means increased CPU utilization.

Other Uses of Virtual Memory

- Some other reasons for having VM are:
 - Running large applications whose address space exceeds the amount of main memory.
 - Decrease apparent startup time for large applications, by allowing applications to start with only a fraction of data resident.
 - Memory-mapped files provide an useful alternative to traditional I/O system calls.

OS and HW support for Virtual Memory

- Virtual address space is partitioned into fixed-size pages (e.g. 4KB).
- Physical memory is partitioned into page frames.
- OS manages **page tables** that define a mapping from (virtual) pages to (physical) page frames.
- HW (MMU) performs address translation on every memory reference.

The OS View of a Process

- Q: So, what is a process, to the OS?
- A: It's just a collection of bookkeeping data, including:
 - CPU register contents to be loaded when process runs.
 - Run state of the process (running, waiting, ready).
 - Memory allocated to the process.
 - Address space structures (e.g., code, data, stack sizes and locations).
 - Other resources in use by the process; such as data describing open files, network connections, etc.
- Linux: this data is managed/accessed by the “task” struct.

Context Switching

- The act of changing between running processes is called a *context switch*.
 - The previous process' context must be saved.
 - The next process' context must be restored.
- Once context save/restore has occurred, the OS can transfer control to the new process.

Processes vs. Threads


- The term **process** usually refers to a single “thread of control” that executes in its own private address space, separate from the address spaces of other processes.
 - Context switches between processes require changing the entire address space, and thus are fairly expensive.
 - Communication between processes requires additional special support from the operating system.

Processes vs. Threads (cont.)

- Modern operating systems support **multithreaded** processes:
 - Multiple cooperating “threads of control” execute within a single process.
 - Threads share most of the process’ context, but have some private data (e.g. their stacks).
 - A full context switch is not required for switching between threads (so threads are “*lightweight*”).


Threads

- A process (with a single address space) may have multiple threads.
- Threads share most of the process' context (e.g. heap, mmap).
- But threads have some private data (e.g. their stacks).



Threads (cont.)

- Threads are concurrent flows of execution belonging to the same program *sharing the same address space*



Next lecture

- Process management