

# Interrupt

*Dongyoon Lee*

# Summary of last lectures

- Tools: building, exploring, and debugging Linux kernel
- Core kernel infrastructure
  - syscall, module, kernel data structures
- Process management & scheduling

# Today: “interrupt”

- A mechanism to implement abstraction and multiplexing
- Interrupt: asking for a service to the kernel
  - by software (e.g., `int`) or by hardware (e.g., keyboard)
- Interrupt handling in Linux
  - top half + bottom half

# Exceptions and Interrupts

- Synchronous interrupts (called exceptions )
  - produced by the CPU while executing instructions
    - non-maskable interrupt (NMI)
- asynchronous interrupts (called interrupts )
  - issued by other hardware devices (e.g., HDD, NIC, etc.)
    - normal and maskable interrupts

# Three kinds of Exceptions


- **Faults** : Can be corrected. The program may continue as if nothing happened (e.g., page faults).
- **Traps** : Reported after the execution of a trapping instruction.
  - invalid memory access, division by zero
  - switch to the kernel mode in response to a system call
  - programmed exceptions: e.g., int 0x03 (breakpoint)
  - If the kernel code does cause a trap, it immediately panics.
- **Aborts** : Some severe unrecoverable error.

# Exceptions

#	Description	Type
0	Divide-by-zero Error	Fault
1	Debug	Fault/Trap
2	Non-maskable Interrupt	Interrupt
3	Breakpoint	Trap
8	Double Fault	Abort
13	General Protection Fault	Fault
14	Page Fault	Fault
...	...	...

**Yeah! New hard disk drive!**

# HDD architecture





# How fast is my new HDD?

- HDD access time = seek time + rotational latency
- Seek time
  - The time to move the disk head to the track that contains data
  - Average seek time: 4 ~ 10 msec
- Rotational latency
  - The delay for the rotation of the disk to bring the required disk sector under the read-write mechanism
  - 7200 RPM: 4.16 msec
- **Access time of your new HDD: about 10 msec**

# Interrupt

- **Compared the the CPU, devices are slow to respond (e.g., 10 msec)**
  - The kernel must be free to go and handle other work, dealing with the hardware only after that hardware has completed its work
- **How to know the completion of a hardware operation**
  - **Polling:** the kernel periodically checks the status of hardware
  - **Interrupt:** the hardware signals its completion to the processor
- Interrupt examples
  - Completion of disk read
  - Key press on a keyboard, network packet arrival

# Interrupt controller



- Interrupts are electrical signals multiplexed by the interrupt controller
  - Sent on a specific pin of the CPU
- Once an interrupt is received, a dedicated function is executed
  - **Interrupt handler**
- The kernel/user space can be interrupted at (nearly) any time to process an interrupt

# Advanced PIC (APIC, I/O APIC)

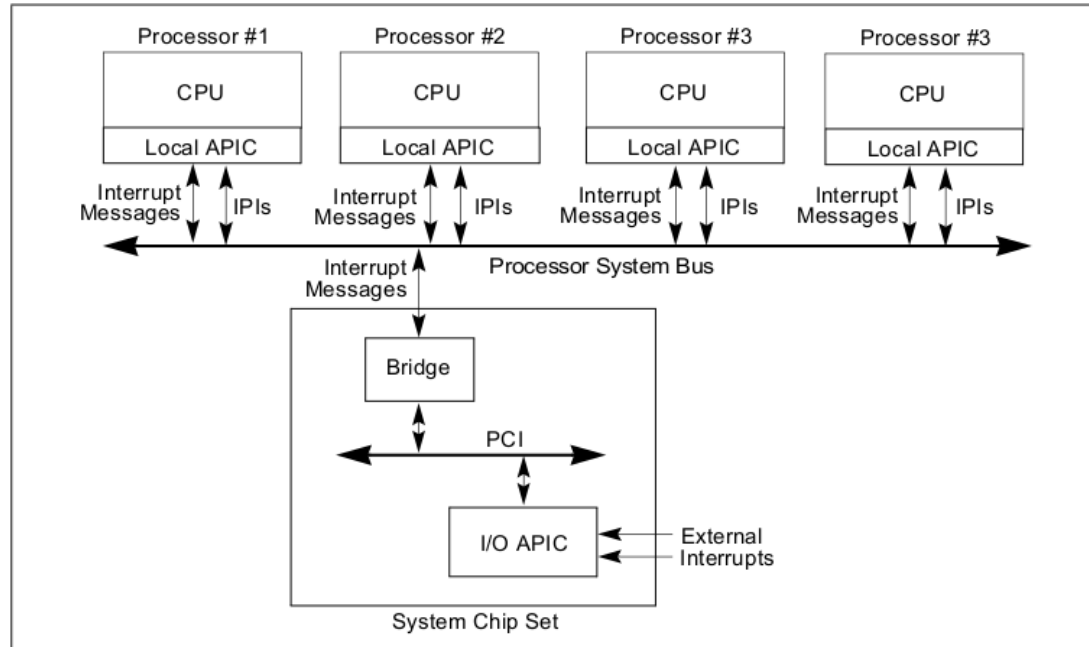
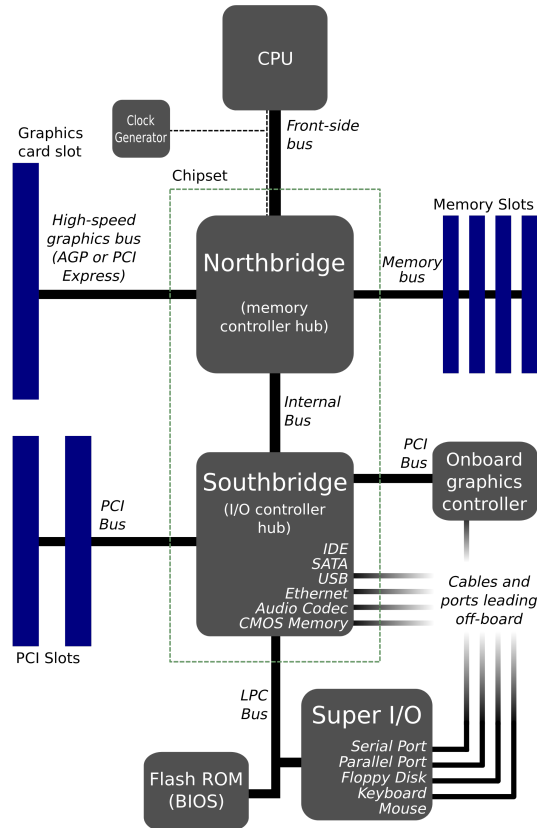


Figure 10-2. Local APICs and I/O APIC When Intel Xeon Processors Are Used in Multiple-Processor Systems


# Advanced PIC (APIC, I/O APIC)

- I/O APIC
  - system chipset (or south bridge)
  - redistribute interrupts to local APICs
- Local APIC
  - inside a processor chip
  - has a timer, which raises timer interrupt
  - issues a IPI (inter-process interrupt)

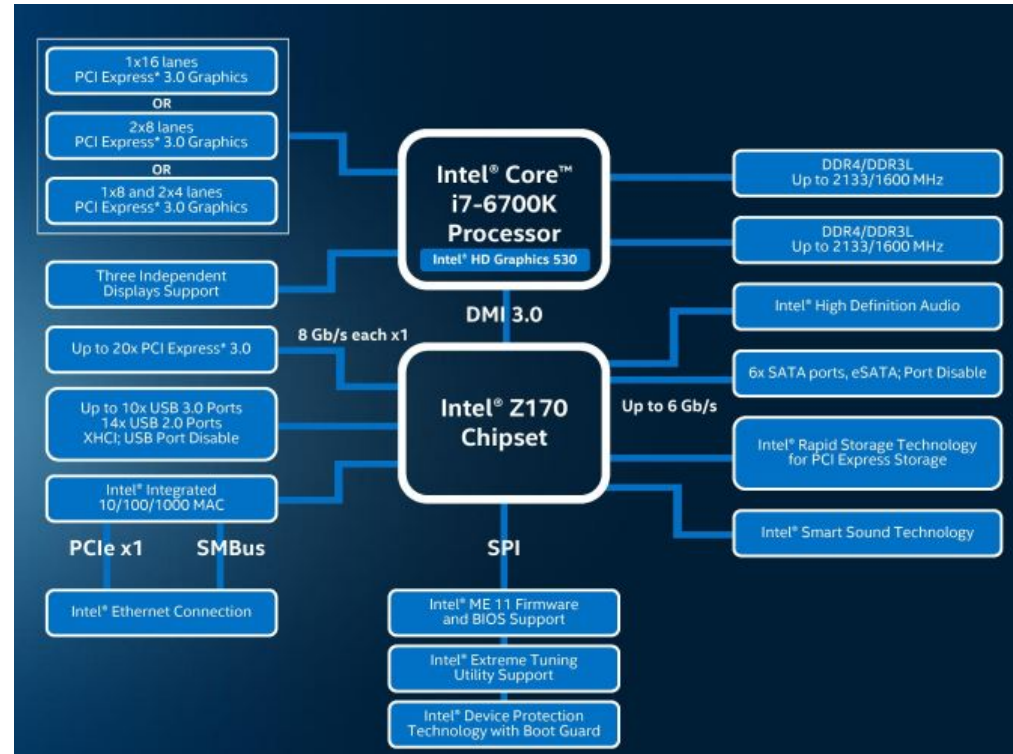
# A bigger picture



# A real motherboard (T42)



# Intel Z170 chipset



- Ref: [The Intel 6th Gen Skylake Review](#)



# Interrupt request (IRQ)

- **Interrupt line** or **interrupt request (IRQ)**
  - device identifier
- E.g., 8259A interrupt lines
  - IRQ 0: system timer, IRQ 1: keyboard controller
  - IRQ 3, 4: serial port, IRQ 5: terminal
- Some interrupt lines can be shared among several devices
  - True for most modern devices (PCIe)

# Next lecture

- Interrupt handler: top half
- Interrupt handler: bottom half