

# Synchronization

CSE 306 Operating Systems  
Dongyoon Lee

# A Simple (Sequential) Queue


```
00: void enqueue (node *new_element) {
01:     node *p;
02:
03:     for (p = head; p->next != NULL; p = p->next) ;
04:     p->next = new_element;
05:     new_element->next = null;
06:
07: }
08:
09: node *dequeue () {
10:     node *element = null;
11:
12:     if (head ->next != null) {
13:         element = head->next;
14:         head->next = head->next->next;
15:     }
16:
17:     return element;
18: }
```

# Lock-based Concurrent Queue

```
00: void enqueue (node *new_element) {
01:     node *p;
02:     LOCK(L);
03:     for (p = head; p->next != NULL; p = p->next) ;
04:     p->next = new_element;
05:     new_element->next = null;
06:     UNLOCK(L);
07: }
08:
09: node *dequeue () {
10:     node *element = null;
11:     LOCK(L);
12:     if (head ->next != null) {
13:         element = head->next;
14:         head->next = head->next->next;
15:     }
16:     UNLOCK(L);
17:     return element;
18: }
```


# What could go wrong?

```
00: void enqueue (node *new_element) {
01:     node *p;
02:     LOCK(L);
03:     for (p = head; p->next != NULL; p = p->next) ;
04:     p->next = new_element;
05:     new_element->next = null;
06:     UNLOCK(L);
07: }
08:
09: node *dequeue () {
10:     node *element = null;
11:     LOCK(L);
12:     if (head ->next != null) {
13:         element = head->next;
14:         head->next = head->next->next;
15:     }
16:     UNLOCK(L);
17:     return element;
18: }
```



# What could go wrong?


```
00: void enqueue (node *new_element) {
01:     node *p;
02:     LOCK(L);
03:     for (p = head; p->next != NULL; p = p->next) ;
04:     p->next = new_element;
05:     new_element->next = null;
06:     UNLOCK(L);
07: }
08:
09: node *dequeue () {
10:     node *element = null;
11:     LOCK(L);
12:     if (head ->next != null) {
13:         element = head->next;
14:         head->next = head->next->next;
15:     }
16:     UNLOCK(L);
17:     return element;
18: }
```



**Two enq()s may find the same tail concurrently.**


# What could go wrong?

```
00: void enqueue (node *new_element) {
01:     node *p;
02:     LOCK(L);
03:     for (p = head; p->next != NULL; p = p->next) ;
04:     p->next = new_element;
05:     new_element->next = null;
06:     UNLOCK(L);
07: }
08:
09: node *dequeue () {
10:     node *element = null;
11:     LOCK(L);
12:     if (head ->next != null) {
13:         element = head->next;
14:         head->next = head->next->next;
15:     }
16:     UNLOCK(L);
17:     return element;
18: }
```



# What could go wrong?

```
00: void enqueue (node *new_element) {
01:     node *p;
02:     LOCK(L);
03:     for (p = head; p->next != NULL; p = p->next) ;
04:     p->next = new_element;
05:     new_element->next = null;
06:     UNLOCK(L);
07: }
08:
09: node *dequeue () {
10:     node *element = null;
11:     LOCK(L);
12:     if (head ->next != null) {
13:         element = head->next;
14:         head->next = head->next->next;
15:     }
16:     UNLOCK(L);
17:     return element;
18: }
```



# Can we make dequeue wait, if empty?


```
00: void enqueue (node *new_element) {
01:     node *p;
02:     LOCK(L);
03:     for (p = head; p->next != NULL; p = p->next) ;
04:     p->next = new_element;
05:     new_element->next = null;
06:     UNLOCK(L);
07: }
08:
09: node *dequeue () {
10:     node *element = null;
11:     LOCK(L);
12:     if (head ->next != null) {
13:         element = head->next;
14:         head->next = head->next->next;
15:     }
16:     UNLOCK(L);
17:     return element;
18: }
```



# Can we make dequeue wait, if empty?

```
00: void enqueue (node *new_element) {
01:     node *p;
02:     LOCK(L);
03:     for (p = head; p->next != NULL; p = p->next) ;
04:     p->next = new_element;
05:     new_element->next = null;
06:     UNLOCK(L);
07: }
08:
09: node *dequeue () {
10:     node *element = null;
11:     LOCK(L);
12:     if (head ->next != null) {
13:         element = head->next;
14:         head->next = head->next->next;
15:     }
16:     UNLOCK(L);
17:     return element;
18: }
```

Add “while”-based waiting



# What could go wrong?

```
00: void enqueue (node *new_element) {
01:     node *p;
02:     LOCK(L);
03:     for (p = head; p->next != NULL; p = p->next) ;
04:     p->next = new_element;
05:     new_element->next = null;
06:     UNLOCK(L);
07: }
08:
09: node *dequeue () {
10:     node *element = null;
11:     LOCK(L);
12:     while (head ->next == null) {}
13:
14:     element = head->next;
15:     head->next = head->next->next;
16:     UNLOCK(L);
17:     return element;
18: }
```

# What could go wrong?


```
00: void enqueue (node *new_element) {  
01:     node *p;  
02:     LOCK(L);  
03:     for (p = head; p->next != NULL; p = p->next) ;  
04:     p->next = new_element;  
05:     new_element->next = null;  
06:     UNLOCK(L);  
07: }
```

```
08:  
09: node *dequeue () {  
10:     node *element = null;  
11:     LOCK(L);  
12:     while (head ->next == null) {}  
13:  
14:     element = head->next;  
15:     head->next = head->next->next;  
16:     UNLOCK(L);  
17:     return element;  
18: }
```

**Problem: Deadlock (hold Lock L forever)**

# What if we acquire Lock after waiting?


```
00: void enqueue (node *new_element) {
01:     node *p;
02:     LOCK(L);
03:     for (p = head; p->next != NULL; p = p->next) ;
04:     p->next = new_element;
05:     new_element->next = null;
06:     UNLOCK(L);
07: }
08:
09: node *dequeue () {
10:     node *element = null;
11:     LOCK(L);
12:     while (head ->next == null) {}
13:
14:     element = head->next;
15:     head->next = head->next->next;
16:     UNLOCK(L);
17:     return element;
18: }
```



# What if we acquire Lock after waiting?

```
00: void enqueue (node *new_element) {
01:     node *p;
02:     LOCK(L);
03:     for (p = head; p->next != NULL; p = p->next) ;
04:     p->next = new_element;
05:     new_element->next = null;
06:     UNLOCK(L);
07: }
08:
09: node *dequeue () {
10:     node *element = null;
11:
12:     while (head ->next == null) {}
13:     LOCK(L);
14:     element = head->next;
15:     head->next = head->next->next;
16:     UNLOCK(L);
17:     return element;
18: }
```


Concurrent deq(), deq(), enq(A)



# What if we acquire Lock after waiting?

```
00: void enqueue (node *new_element) {
01:     node *p;
02:     LOCK(L);
03:     for (p = head; p->next != NULL; p = p->next) ;
04:     p->next = new_element;
05:     new_element->next = null;
06:     UNLOCK(L);
07: }
08:
09: node *dequeue () {
10:     node *element = null;
11:
12:     while (head ->next == null) {}
13:     LOCK(L);
14:     element = head->next;
15:     head->next = head->next->next;
16:     UNLOCK(L);
17:     return element;
18: }
```


Concurrent deq(), deq(), enq(A)



# What if we acquire Lock after waiting?

```
00: void enqueue (node *new_element) {
01:     node *p;
02:     LOCK(L);
03:     for (p = head; p->next != NULL; p = p->next) ;
04:     p->next = new_element;
05:     new_element->next = null;
06:     UNLOCK(L);
07: }
08:
09: node *dequeue () {
10:     node *element = null;
11:
12:     while (head ->next == null) {}
13:     LOCK(L);
14:     element = head->next;
15:     head->next = head->next->next;
16:     UNLOCK(L);
17:     return element;
18: }
```

Concurrent deq(), deq(), enq(A)



Two deq()s may find the queue not-empty concurrently.

# A possible solution. Another problem?

```
00: void enqueue (node *new_element) {
01:     node *p;
02:     LOCK(L);
03:     for (p = head; p->next != NULL; p = p->next) ;
04:     p->next = new_element;
05:     new_element->next = null;
06:     UNLOCK(L);
07: }
08:
09: node *dequeue () {
10:     node *element = null;
11:     LOCK(L);
12:     while (head ->next == null) {UNLOCK(L); LOCK(L);}
13:
14:     element = head->next;
15:     head->next = head->next->next;
16:     UNLOCK(L);
17:     return element;
18: }
```

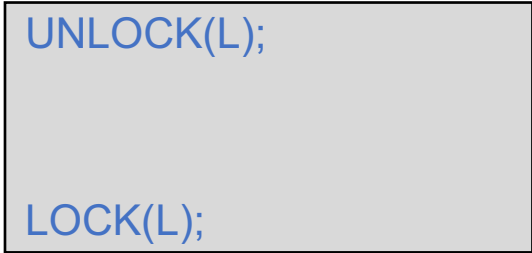


# A possible solution. Another problem?

```
00: void enqueue (node *new_element) {
01:     node *p;
02:     LOCK(L);
03:     for (p = head; p->next != NULL; p = p->next) ;
04:     p->next = new_element;
05:     new_element->next = null;
06:     UNLOCK(L);
07: }
08:
09: node *dequeue () {
10:     node *element = null;
11:     LOCK(L);
12:     while (head ->next == null) {UNLOCK(L); LOCK(L);}
13:
14:     element = head->next;
15:     head->next = head->next->next;
16:     UNLOCK(L);
17:     return element;
18: }
```

**Problem: Busy waiting.**

# How to avoid busy waiting?

```
00: node *dequeue () {  
01:     node *element = null;  
02:     LOCK(L);  
03:     while (head ->next == null) {  
04:         UNLOCK(L);  
05:           
06:         LOCK(L);  
07:     }  
08:     element = head->next;  
09:     head->next = head->next->next;  
10:     UNLOCK(L);  
11:     return element;  
12: }  
13: }
```

# How to avoid busy waiting?

```
00: node *dequeue () {
01:     node *element = null;
02:     LOCK(L);
03:     while (head ->next == null) {
04:         UNLOCK(L);
05:         add myself into a waitlist;
06:         go to sleep;
07:         LOCK(L);
08:     }
09:     element = head->next;
10:     head->next = head->next->next;
11:     UNLOCK(L);
12:     return element;
13: }
```

**WAIT (L, CV)**  
- **Atomically by OS**

# How to avoid busy waiting?

```
00: node *dequeue () {  
01:     node *element = null;  
02:     LOCK(L);  
03:     while (head ->next == null) {  
04:         UNLOCK(L);  
05:         add myself into a waitlist;  
06:         go to sleep;  
07:         LOCK(L);  
08:     }  
09:     element = head->next;  
10:     head->next = head->next->next;  
11:     UNLOCK(L);  
12:     return element;  
13: }
```

## WAIT (L, CV)

- Atomically by OS
- If not,
  - T1: put into a waitlist. before go to sleep
  - T2: wake up and send signal
  - T1: go to sleep
  - We “lost the wake-up signal”.

# How to avoid busy waiting?

```
00: void enqueue (node *new_element) {  
01:     node *p;  
02:     LOCK(L);  
03:     for (p = head; p->next != NULL; p = p->next) ;  
04:     p->next = new_element;  
05:     new_element->next = null;  
06:     take a waiter off the waitlist;  
07:     wake up;  
08:     UNLOCK(L);  
09: }  
10:
```

# How to avoid busy waiting?

00: void enqueue (node \*new\_element) {

01: node \*p;

02: LOCK(L);

03: for (p = head; p->next != NULL; p = p->next) ;

04: p->next = new\_element;

05: new\_element->next = null;

06: take a waiter off the waitlist;

07: wake up;

08: UNLOCK(L);

09: }

10:


SIGNAL (CV) or BROADCAST (CV)

# Locks and Conditional Variables


```
00: void enqueue (node *new_element) {
01:     node *p;
02:     LOCK(L);
03:     for (p = head; p->next != NULL; p = p->next) ;
04:     p->next = new_element;
05:     new_element->next = null;
06:     SIGNAL(CV);
07:     UNLOCK(L);
08: }
09: node *dequeue () {
10:     node *element = null;
11:     LOCK(L);
12:     while (head ->next == null) {
13:         WAIT(L, CV);
14:     }
15:     element = head->next;
16:     head->next = head->next->next;
17:     UNLOCK(L);
18:     return element;
19: }
```

# Synchronization: Don't do simultaneously!

- Locks: Mutual exclusion

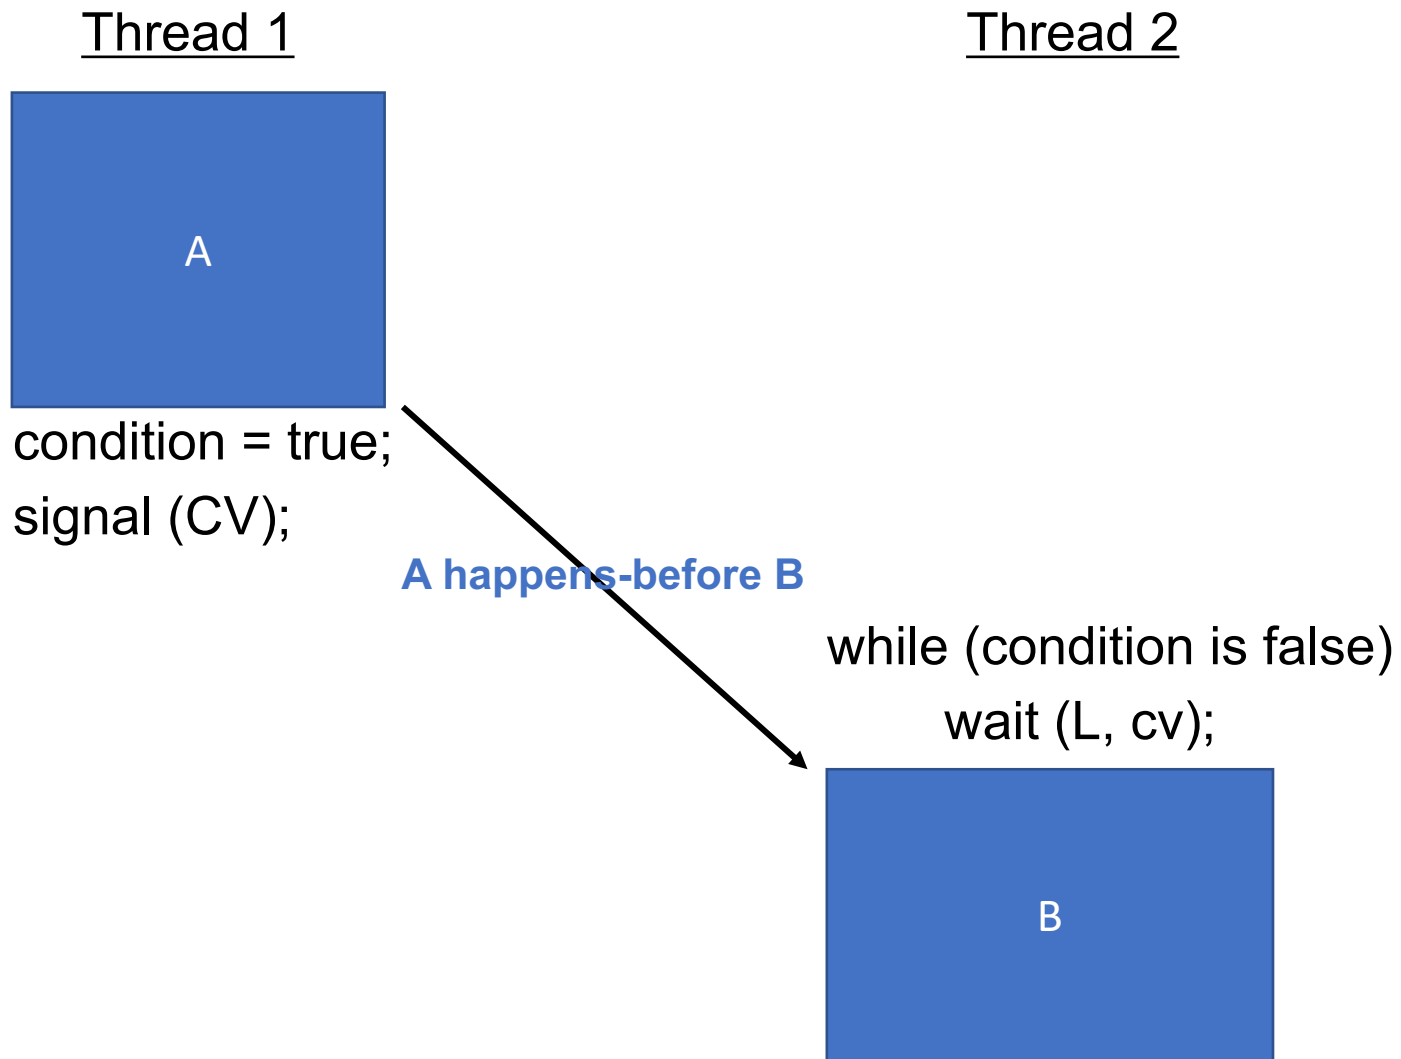


- Locks + CVs: Ordering constraint







# Monitor: 1 lock and 0~n conditional vars



# Monitor: 1 lock and 0~n conditional vars



# Monitor: 1 lock and 0~n conditional vars



# How to build a monitor (step by step)

1. Identify shared data
2. Assign locks to shared data
3. Add **lock** (before r/w) and **unlock** (after r/w)
4. List before and after condition (e.g., A has to happen before B) and assign a conditional variable to each one
5. Add **wait** and while loop
6. Add **signal** or **broadcast** (whenever a condition is changed)
7. Make sure reestablish program invariants

# Coke Machine Producer and Consumer

```
int numCoke;
```

```
Queue
```

```
MAX_COKES
```

```
00: Consumer () {
```

```
01:
```

```
02:
```

```
03:
```

```
04:
```

```
05:     takeCokeOutofMachine();
```

```
06:     numCoke--;
```

```
07:
```

```
08:
```

```
09: }
```

```
10: Producer () {
```

```
11:
```

```
12:
```

```
13:
```

```
14:
```

```
15:     putCokeIntoMachine();
```

```
16:     numCoke++;
```

```
17:
```

```
18:
```

```
19: }
```

# Coke Machine Producer and Consumer

## 1. Identify shared data

int numCoke;

Queue

MAX\_COKES

## 2. Assign locks: CokeLock

```
00: Consumer () {  
01:  
02:  
03:  
04:  
05:     takeCokeOutofMachine();  
06:     numCoke--;  
07:  
08:  
09: }
```

```
10: Producer () {  
11:  
12:  
13:  
14:  
15:     putCokeIntoMachine();  
16:     numCoke++;  
17:  
18:  
19: }
```

# Coke Machine Producer and Consumer

## 1. Identify shared data

int numCoke;

Queue

MAX\_COKES

## 2. Assign locks: CokeLock

## 3. Add lock/unlock

```
00: Consumer () {  
01:     LOCK(CokeLock);  
02:  
03:  
04:  
05:     takeCokeOutofMachine();  
06:     numCoke--;  
07:  
08:     UNLOCK(CokeLock);  
09: }
```

```
10: Producer () {  
11:     LOCK(CokeLock);  
12:  
13:  
14:  
15:     putCokeIntoMachine();  
16:     numCoke++;  
17:  
18:     UNLOCK(CokeLock);  
19: }
```

# Coke Machine Producer and Consumer

## 4. List before/after conditions and assign conditional variables

int numCoke;

Queue

MAX\_COKES

A) Must be >0 cokes

before we take a Coke out

B) Must be < MAX\_COKES cokes

before we put a Coke in

```
00: Consumer () {  
01:     LOCK(CokeLock);  
02:  
03:  
04:  
05:     takeCokeOutofMachine();  
06:     numCoke--;  
07:  
08:     UNLOCK(CokeLock);  
09: }
```

```
10: Producer () {  
11:     LOCK(CokeLock);  
12:  
13:  
14:  
15:     putCokeIntoMachine();  
16:     numCoke++;  
17:  
18:     UNLOCK(CokeLock);  
19: }
```



# Coke Machine Producer and Consumer

## 4. List before/after conditions and assign conditional variables

```
int numCoke;  
Queue  
MAX_COKES
```

**noCoke** A) Must be >0 cokes  
before we take a Coke out

**noSpace** B) Must be < MAX\_COKES cokes  
before we put a Coke in

```
00: Consumer () {  
01:     LOCK(CokeLock);  
02:  
03:  
04:  
05:     takeCokeOutofMachine();  
06:     numCoke--;  
07:  
08:     UNLOCK(CokeLock);  
09: }
```

```
10: Producer () {  
11:     LOCK(CokeLock);  
12:  
13:  
14:  
15:     putCokeIntoMachine();  
16:     numCoke++;  
17:  
18:     UNLOCK(CokeLock);  
19: }
```

# Coke Machine Producer and Consumer

int numCokes;

Queue

MAX\_COKES

A) Must be >0 cokes

before we take a Coke out

B) Must be < MAX\_COKES cokes

before we put a Coke in

## 5. Add wait and while loop

```
00: Consumer () {  
01:     LOCK (CokeLock);  
02:     while (numCoke == 0) {  
03:         WAIT(CokeLock, noCoke);  
04:     }  
05:     takeCokeOutofMachine();  
06:     numCoke--;  
07:  
08:     UNLOCK(CokeLock);  
09: }
```

```
10: Producer () {  
11:     LOCK (CokeLock);  
12:     while (numCoke == MAX_COKES) {  
13:         WAIT(CokeLock, noSpace);  
14:     }  
15:     putCokeIntoMachine();  
16:     numCoke++;  
17:  
18:     UNLOCK(CokeLock);  
19: }
```

# Coke Machine Producer and Consumer

int numCokes;

Queue

MAX\_COKES

A) Must be >0 cokes

before we take a Coke out

B) Must be < MAX\_COKES cokes

before we put a Coke in

```
00: Consumer () {
01:     LOCK (CokeLock);
02:     while (numCoke == 0) {
03:         WAIT(CokeLock, noCoke);
04:     }
05:     takeCokeOutofMachine();
06:     numCoke--;
07:     SIGNAL(noSpace); 6. Add signal
08:     UNLOCK(CokeLock);
09: }
```

```
10: Producer () {
11:     LOCK (CokeLock);
12:     while (numCoke == MAX_COKES) {
13:         WAIT(CokeLock, noSpace);
14:     }
15:     putCokeIntoMachine();
16:     numCoke++;
17:     SIGNAL(noCoke);
18:     UNLOCK(CokeLock);
19: }
```

# Semaphore

- Combine two abstractions (lock and conditional variables)
- Non-negative integer value
- Two operations

```
DOWN () {  
    do{  
        if (value > 0){  
            value--;  
            break;  
        }  
    } while (1);  
}  
  
UP() {  
    value++;  
}
```


**Atomic**

# Semaphore

- **Critical Section**

- `sem_init(1)`: initial value is 1.
- DOWN: wait if 0, otherwise set 0.  $\sim$  = LOCK
- UP: set 1  $\sim$  = UNLOCK


DOWN(sem);



UP(sem);

- **Ordering**


- `sem_init(0)`: initial value is 0.
- DOWN: wait if 0, otherwise set 0.  $\sim$  = WAIT
- UP: set 1  $\sim$  = SIGNAL



UP(sem);

**A happens-  
before B**

DOWN(sem);



# Coke Machine Semaphore

```
sem Mutex = sem_init (1);
```

```
00: Consumer () {  
01:  
02:     down (Mutex);  
03:     takeCokeOutofMachine();  
04:     numCoke--;  
05:     up (Mutex);  
06:  
07: }  
08:  
09:
```

```
10: Producer () {  
11:  
12:     down (Mutex);  
13:     putCokeIntoMachine();  
14:     numCoke++;  
15:     up (Mutex);  
16:  
17: }  
18:  
19:
```

# Coke Machine Semaphore

```
sem Mutex = sem_init (1);  
sem Full = sem_init (0);
```

```
00: Consumer () {  
01:     down (Full);  
02:     down (Mutex);  
03:     takeCokeOutofMachine();  
04:     numCoke--;  
05:     up (Mutex);  
06:  
07: }  
08:  
09:
```

```
10: Producer () {  
11:  
12:     down (Mutex);  
13:     putCokeIntoMachine();  
14:     numCoke++;  
15:     up (Mutex);  
16:     up (Full);  
17: }  
18:  
19:
```

# Coke Machine Semaphore

```
sem Mutex = sem_init (1);  
sem Full = sem_init (0);  
sem Empty = sem_init (N);
```

```
00: Consumer () {  
01:     down (Full);  
02:     down (Mutex);  
03:     takeCokeOutofMachine();  
04:     numCoke--;  
05:     up (Mutex);  
06:     up (Empty);  
07: }  
08:  
09:
```

```
10: Producer () {  
11:     down (Empty);  
12:     down (Mutex);  
13:     putCokeIntoMachine();  
14:     numCoke++;  
15:     up (Mutex);  
16:     up (Full);  
17: }  
18:  
19:
```