

Process Address Space

Dongyoon Lee

Summary of last lectures

- Tools: building, exploring, and debugging Linux kernel
- Core kernel infrastructure
 - syscall, module, kernel data structures
- Process management & scheduling
- Interrupt & interrupt handler
- Kernel synchronization
- Memory management

Today: process address space

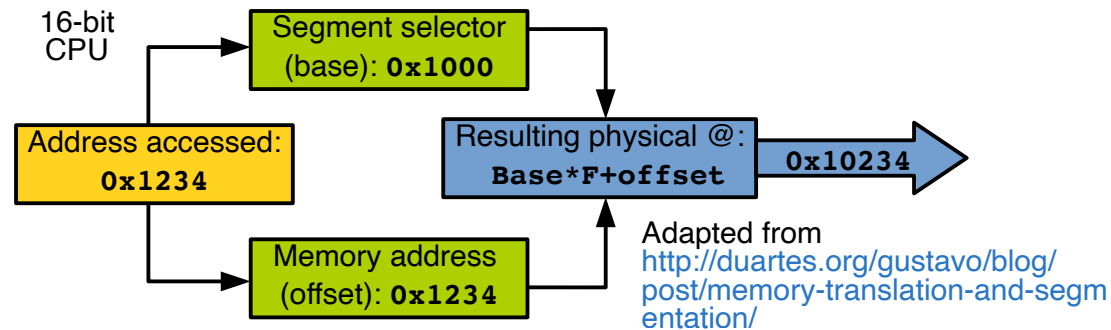
- Address space
- Memory descriptor: `mm_struct`
- Virtual Memory Area (VMA)
- VMA manipulation
- Page tables

Address space

- The memory that a process can access
 - Illusion that the process can access 100% of the system memory
 - With virtual memory, can be much larger than the actual amount of physical memory
- Defined by the process page table set up by the kernel

Address space

- A memory address is an index within the address spaces:
 - Identify a specific byte
- Each process is given a flat 32/64-bits address space
 - Not segmented



Address space

- **Virtual Memory Areas (VMA)**


- Interval of addresses that the process has the right to access
- Can be dynamically added or removed to the process address space
- Associated permissions: read, write, execute
- *Illegal access → segmentation fault*

```
$ cat /proc/1/maps          # or sudo pmap 1
55fe3bf02000-55fe3bff9000 r-xp 00000000 fd:00 1975429 /usr/lib/systemd/systemd
55fe3bfffa000-55fe3c021000 r--p 000f7000 fd:00 1975429 /usr/lib/systemd/systemd
55fe3c021000-55fe3c022000 rw-p 0011e000 fd:00 1975429 /usr/lib/systemd/systemd
55fe3db4a000-55fe3ddfd000 rw-p 00000000 00:00 0 [heap]
7f7522769000-7f7522fd9000 rw-p 00000000 00:00 0
7f7523150000-7f7523265000 r-xp 00000000 fd:00 1979800 /usr/lib64/libm-2.25.so
7f7523265000-7f7523464000 ---p 00115000 fd:00 1979800 /usr/lib64/libm-2.25.so
7f7523464000-7f7523465000 r--p 00114000 fd:00 1979800 /usr/lib64/libm-2.25.so
7f7523465000-7f7523466000 rw-p 00115000 fd:00 1979800 /usr/lib64/libm-2.25.so
```

Address space

- **VMAs can contain:**
 - Mapping of the executable file code (*text section*)
 - Mapping of the executable file initialized variables (*data section*)
 - Mapping of the zero page for uninitialized variables (*bss section*)
 - Mapping of the zero page for the *user-space stack*
 - Text, data, bss for each *shared library* used
 - Memory-mapped files, shared memory segment, anonymous mappings (used by malloc)

Address space



Memory descriptor: `mm_struct`

- Address space in linux kernel: `struct mm_struct`

```
/* linux/include/linux/mm types.h */
```

```
struct mm_struct {
    struct vm_area_struct *mmap;           /* list of VMAs */
    struct rb_root mm_rb;                 /* rbtree of VMAs */
    pgd_t *pgd;                           /* page global directory */
    atomic_t mm_users;                    /* address space users */
    atomic_t mm_count;                    /* primary usage counters */
    int map_count;                         /* number of VMAs */
    struct rw_semaphore mmap_sem;         /* VMA semaphore */
    spinlock_t page_table_lock;           /* page table lock */
    struct list_head mmlist;              /* list of all mm_struct */
    unsigned long start_code;              /* start address of code */
    unsigned long end_code;                /* end address of code */
    unsigned long start_data;              /* start address of data */
    unsigned long end_data;                /* end address of data */
    unsigned long start_brk;               /* start address of heap */
    unsigned long end_brk;                 /* end address of heap */
    unsigned long start_stack;             /* start address of stack */
    /* ... */
}
```

Memory descriptor: `mm_struct`

```
unsigned long    arg_start;    /* start of arguments */
unsigned long    arg_end;      /* end of arguments */
unsigned long    env_start;    /* start of environment */
unsigned long    total_vm;     /* total pages mapped */
unsigned long    locked_vm;    /* number of locked pages */
unsigned long    flags;        /* architecture specific data */
spinlock_t       ioctx_lock;   /* Asynchronous I/O list lock */
/* ... */
};
```

- `mm_users` : number of processes (threads) using the address space
- `mm_count` : reference count:
 - +1 if `mm_users` > 0
 - +1 if the kernel is using the address space

Memory descriptor: `mm_struct`

- `mmap` and `mm_rb` are respectively a linked list and a tree containing all the VMAs in this address space
 - List used to iterate over all the VMAs in an ascending order
 - Tree used to find a specific VMA
- All `mm_struct` are linked together in a doubly linked list
 - Through the `mmlist` field of the `mm_struct`

Allocating a memory descriptor

- A task memory descriptor is located in the `mm` field of the corresponding `task_struct`

```
/* linux/include/linux/sched.h */

struct task_struct {
    struct thread_info    thread_info;
    /* ... */
    const struct sched_class *sched_class;
    struct sched_entity   se;
    struct sched_rt_entity rt;
    /* ... */
    struct mm_struct      *mm;
    struct mm_struct      *active_mm;
    /* ... */
};
```

Allocating a memory descriptor

- Current task memory descriptor: `current->mm`
- During `fork()`, `copy_mm()` is making a copy of the parent memory descriptor for the child
 - `copy_mm()` calls `dup_mm()` which calls `allocate_mm()` → allocates a `mm` struct object from a slab cache
- Two threads sharing the same address space have the `mm` field of their `task_struct` pointing to the same `mm_struct` object
 - Threads are created using the `CLONE_VM` flag passed to `clone()` → `allocate_mm()` is not called

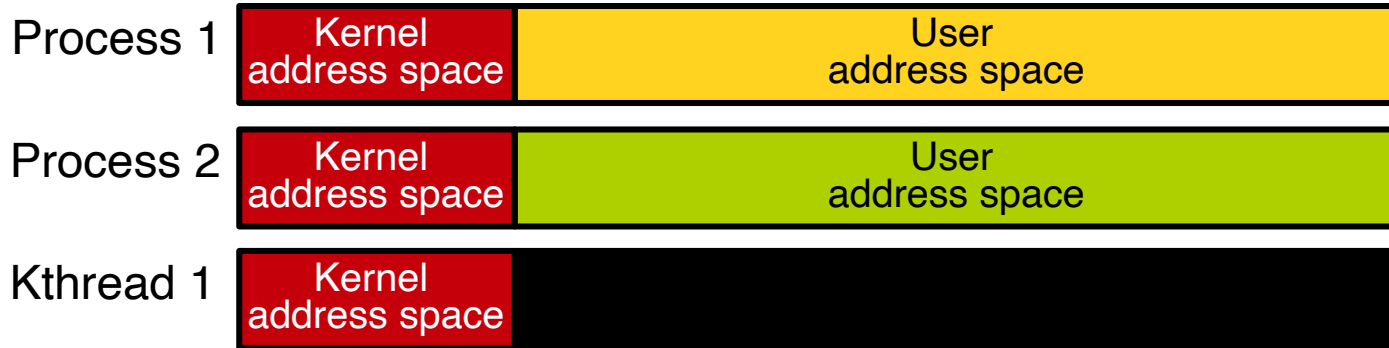
Destroying a memory descriptor

- When a process exits, `do_exit()` is called and it calls `exit_mm()`
 - Performs some housekeeping/statistics updates and calls `mmapput()`

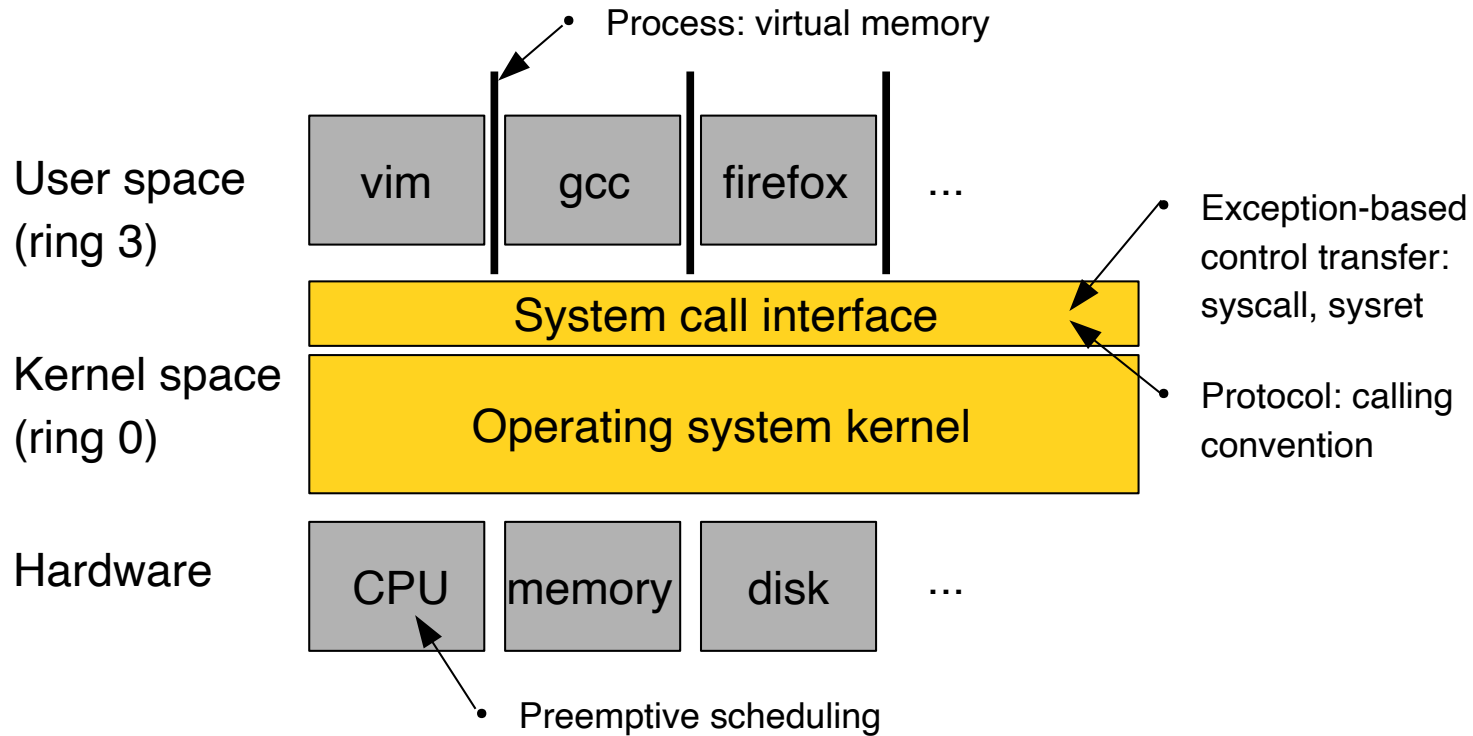
```
void mmapput(struct mm_struct *mm) {
    might_sleep();
    if (atomic_dec_and_test(&mm->mm_users))
        __mmapput(mm);
}
static inline void __mmapput(struct mm_struct *mm) {
    /* ... */
    mmdrop(mm);
}
static inline void mmdrop(struct mm_struct *mm) {
    if (unlikely(atomic_dec_and_test(&mm->mm_count)))
        __mmdrop(mm);
}
void __mmdrop(struct mm_struct *mm) {
    /* ... */
    free_mm(mm);
}
```

The `mm_struct` and kernel threads

- Kernel threads do not have a user-space address space
 - `mm` field of a kernel thread `task_struct` is `NULL`



The `mm_struct` and kernel threads



The `mm_struct` and kernel threads

- However kernel threads still need to access the kernel address space
 - When a kernel thread is scheduled, the kernel notice its `mm` is `NULL` so it keeps the previous address space loaded (page tables)
 - Kernel makes the `active_mm` field of the kernel thread to point on the borrowed `mm_struct`
 - It is okay because the kernel address space is the same in all tasks

Review: Segmentation in x86

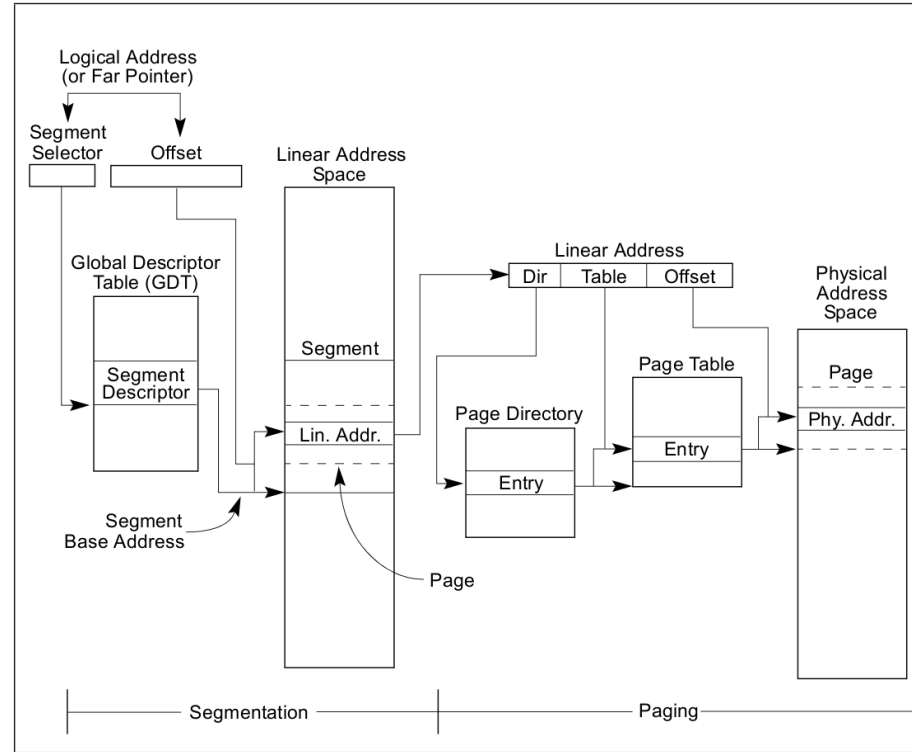


Figure 3-1. Segmentation and Paging

Review: Privilege levels of a segment

- CPL (current privilege level)
 - the privilege level of currently executing program
 - bits 0 and 1 in the `%cs` register
- RPL (requested privilege level)
 - an override privilege level that is assigned to a segment selector
 - a segment selector is a part (16-bit) of segment registers (e.g., `ds`, `fs`), which is an index of a segment descriptor and RPL
- DPL (descriptor privilege level)
 - the privilege level of a segment

Review: How isolation is enforced in x86?




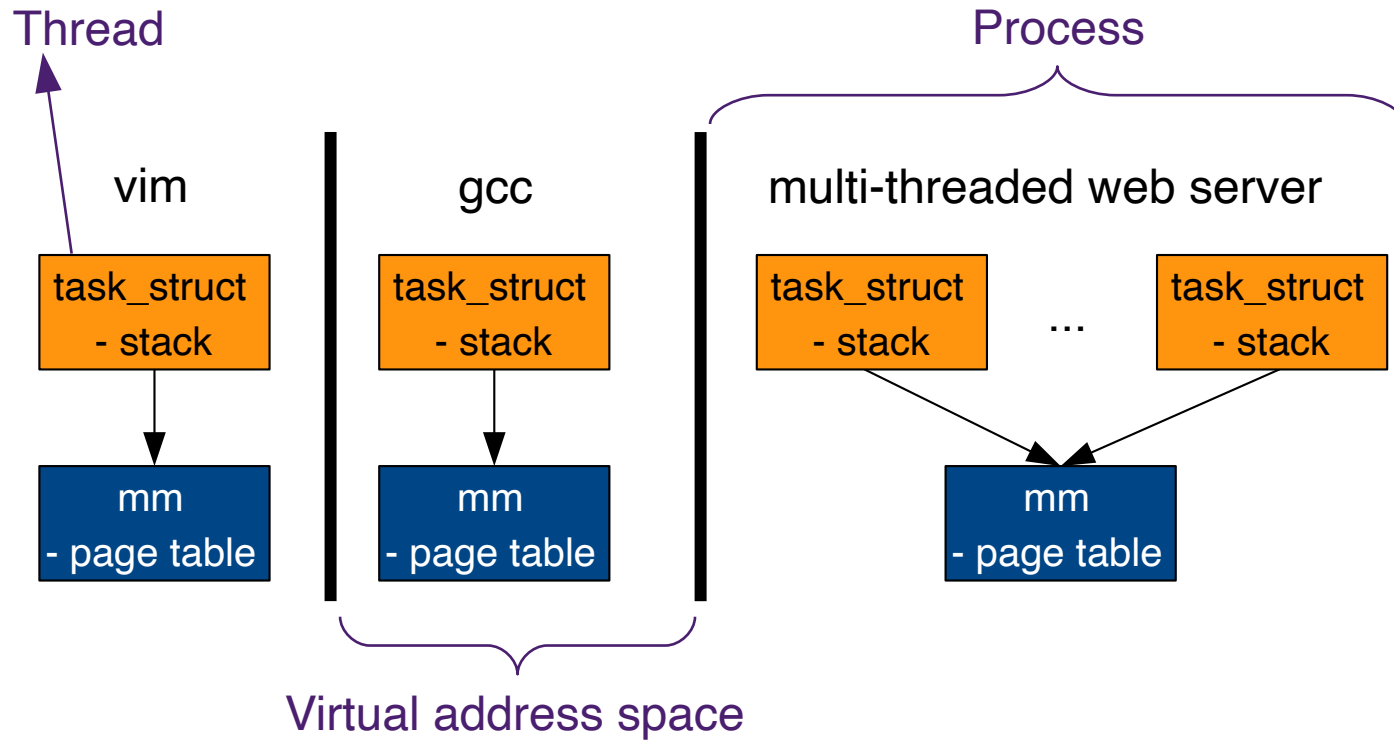
Figure 5-4. Privilege Check for Data Access

- Access is granted if $RPL \geq CPL$ and $DPL \geq CPL$

Review: How to switch b/w rings (ring 0 ↔ ring 3)?

- Controlled transfer: system call
 - `int`, `sysenter` or `syscall` instruction set CPL to 0; change to `KERNEL_CS` and `KERNEL_DS` segments
 - set CPL to 3 before going back to user space; change to `USER_CS` and `USER_DS` segments

mm_struct vs. task_struct



Virtual Memory Area (VMA)

- Each line corresponds to one VMA

```
$ cat /proc/1/maps # or sudo pmap 1
55fe3bf02000-55fe3bff9000 r-xp 00000000 fd:00 1975429 /usr/lib/systemd/systemd
55fe3bffa000-55fe3c021000 r--p 000f7000 fd:00 1975429 /usr/lib/systemd/systemd
55fe3c021000-55fe3c022000 rw-p 0011e000 fd:00 1975429 /usr/lib/systemd/systemd
55fe3db4a000-55fe3ddfd000 rw-p 00000000 00:00 0 [heap]
7f7522769000-7f7522fd9000 rw-p 00000000 00:00 0
7f7523150000-7f7523265000 r-xp 00000000 fd:00 1979800 /usr/lib64/libm-2.25.so
7f7523265000-7f7523464000 ---p 00115000 fd:00 1979800 /usr/lib64/libm-2.25.so
7f7523464000-7f7523465000 r--p 00114000 fd:00 1979800 /usr/lib64/libm-2.25.so
7f7523465000-7f7523466000 rw-p 00115000 fd:00 1979800 /usr/lib64/libm-2.25.so
```

```
# r = read
# w = write
# x = execute
# s = shared
# p = private (copy on write)
```

Virtual Memory Area (VMA)

- Each VMA is represented by an object of type `vm_area_struct`

```
/* linux/include/linux/mm_types.h */
```

```
struct vm_area_struct {
    struct                mm_struct *vm_mm; /* associated address space (mm_struct) */
    unsigned long         vm_start;        /* VMA start, inclusive */
    unsigned long         vm_end;          /* VMA end, exclusive */
    struct vm_area_struct *vm_next;        /* list of VMAs */
    struct vm_area_struct *vm_prev;        /* list of VMAs */
    pgprot_t              vm_page_prot;    /* access permissions */
    unsigned long         vm_flags;        /* flags */
    struct rb_node         vm_rb;           /* VMA node in the tree */
    struct list_head       anon_vma_chain; /* list of anonymous mappings */
    struct anon_vma        *anon_vma;      /* anonmous vma object */
    struct vm_operation_struct *vm_ops;     /* operations */
    unsigned long         vm_pgoff;        /* offset within file */
    struct file            *vm_file;        /* mapped file (can be NULL) */
    void                  *vm_private_data; /* private data */
    /* ... */
}
```


Virtual Memory Area (VMA)

- The VMA exists over `[vm_start, vm_end)` in the corresponding address space → size in bytes: `vm_end - vm_start`
- Address space is pointed by the `vm_mm` field (of type `mm_struct`)
- Each VMA is unique to the associated `mm_struct`
 - Two processes mapping the same file will have two different `mm_struct` objects, and two different `vm_area_struct` objects
 - Two threads sharing a `mm_struct` object also share the `vm_area_struct` objects

VMA flags

Flag	Effect on the VMA and Its Pages
VM_READ	Pages can be read from.
VM_WRITE	Pages can be written to.
VM_EXEC	Pages can be executed.
VM_SHARED	Pages are shared.
VM_MAYREAD	The VM_READ flag can be set.
VM_MAYWRITE	The VM_WRITE flag can be set.
VM_MAYEXEC	The VM_EXEC flag can be set.
VM_MAYSHARE	The VM_SHARE flag can be set.

VMA flags

Flag	Effect on the VMA and Its Pages
VM_GROWSDOWN	The area can grow downward.
VM_GROWSUP	The area can grow upward.
VM_SHM	The area is used for shared memory.
VM_DENYWRITE	The area maps an unwritable file.
VM_EXECUTABLE	The area maps an executable file.
VM_LOCKED	The pages in this area are locked.
VM_IO	The area maps a device's I/O space.
VM_SEQ_READ	The pages seem to be accessed sequentially.

VMA flags

Flag	Effect on the VMA and Its Pages
VM_RAND_READ	The pages seem to be accessed randomly.
VM_DONTCOPY	This area must not be copied on fork().
VM_DONTEXPAND	This area cannot grow via mremap().
VM_RESERVED	This area must not be swapped out.
VM_ACCOUNT	This area is an accounted VM object.
VM_HUGETLB	This area uses hugetlb pages.
VM_NONLINEAR	This area is a nonlinear mapping.

VMA flags

- Combining `VM_READ` , `VM_WRITE` and `VM_EXEC` gives the permissions for the entire area, for example:
 - Object code is `VM_READ` and `VM_EXEC`
 - Stack is `VM_READ` and `VM_WRITE`
- `VM_SEQ_READ` and `VM_RAND_READ` are set through the `madvise()` system call
 - Instructs the file pre-fetching algorithm read-ahead to increase or decrease its pre-fetch window

VMA flags

- `VM_HUGETLB` indicates that the area uses pages larger than the regular size
 - 2M and 1G on x86
 - Larger page size → less TLB miss → faster memory access

VMA operations

- `vm_ops` in `vm_area_struct` is a struct of function pointers to operate on a specific VMA

```
/* linux/include/linux/mm.h */
struct vm_operations_struct {
    /* called when the area is added to an address space */
    void (*open)(struct vm_area_struct * area);

    /* called when the area is removed from an address space */
    void (*close)(struct vm_area_struct * area);

    /* invoked by the page fault handler when a page that is
     * not present in physical memory is accessed*/
    int (*fault)(struct vm_area_struct *vma, struct vm_fault *vmf);

    /* invoked by the page fault handler when a previously read-only
     * page is made writable */
    int (*page_mkwrite)(struct vm_area_struct *vma, struct vm_fault *vmf);
    /* ... */
}
```

VMA manipulation: **find_vma()**

```
/* linux/mm/mmap.c */

/* Look up the first VMA which satisfies  addr < vm_end,  NULL if none. */
struct vm_area_struct *find_vma(struct mm_struct *mm, unsigned long addr)
{
    struct rb_node *rb_node;
    struct vm_area_struct *vma;

    /* Check the cache first. */
    vma = vmacache_find(mm, addr);
    if (likely(vma))
        return vma;

    rb_node = mm->mm_rb.rb_node;

    while (rb_node) {
        struct vm_area_struct *tmp;

        tmp = rb_entry(rb_node, struct vm_area_struct, vm_rb);
```


VMA manipulation: **find_vma()**

```
    if (tmp->vm_end > addr) {  
        vma = tmp;  
        if (tmp->vm_start <= addr)  
            break;  
        rb_node = rb_node->rb_left;  
    } else  
        rb_node = rb_node->rb_right;  
}  
  
if (vma)  
    vmacache_update(addr, vma);  
return vma;  
}
```

VMA manipulation

```
/* linux/include/linux/mm.h */

/* Look up the first VMA which satisfies  addr < vm_end,  NULL if none. */
struct vm_area_struct *find_vma(struct mm_struct *mm, unsigned long addr);

/*
 * Same as find_vma, but also return a pointer to the previous VMA in *pprev.
 */
struct vm_area_struct *
find_vma_prev(struct mm_struct *mm, unsigned long addr,
              struct vm_area_struct **pprev);

/* Look up the first VMA which intersects the interval start_addr..end_addr-1,
   NULL if none.  Assume start_addr < end_addr. */
struct vm_area_struct * find_vma_intersection(struct mm_struct * mm,
                                              unsigned long start_addr, unsigned long end_addr);
```

Creating an address interval

- `do_mmap()` is used to create a new linear address interval:
 - Can result in the creation of a new VMAs
 - Or a merge of the create area with an adjacent one when they have the same permissions

```
/*  
 * The caller must hold down_write(&current->mm->mmap_sem).  
 */  
unsigned long do_mmap(struct file *file, unsigned long addr,  
                     unsigned long len, unsigned long prot,  
                     unsigned long flags, vm_flags_t vm_flags,  
                     unsigned long pgoff, unsigned long *populate,  
                     struct list_head *uf);
```

Creating an address interval

- `prot` specifies access permissions for the memory pages

Flag	Effect on the new interval
<code>PROT_READ</code>	Corresponds to <code>VM_READ</code>
<code>PROT_WRITE</code>	Corresponds to <code>VM_WRITE</code>
<code>PROT_EXEC</code>	Corresponds to <code>VM_EXEC</code>
<code>PROT_NONE</code>	Cannot access page

Creating an address interval

- `flags` specifies the rest of the VMA options

Flag	Effect on the new interval
<code>MAP_SHARED</code>	The mapping can be shared.
<code>MAP_PRIVATE</code>	The mapping cannot be shared.
<code>MAP_FIXED</code>	The new interval must start at the given address addr.
<code>MAP_ANONYMOUS</code>	The mapping is not file-backed, but is anonymous.
<code>MAP_GROWSDOWN</code>	Corresponds to <code>VM_GROWSDOWN</code> .

Creating an address interval

- `flags` specifies the rest of the VMA options

Flag	Effect on the new interval
<code>MAP_DENYWRITE</code>	Corresponds to <code>VM_DENYWRITE</code> .
<code>MAP_EXECUTABLE</code>	Corresponds to <code>VM_EXECUTABLE</code> .
<code>MAP_LOCKED</code>	Corresponds to <code>VM_LOCKED</code> .
<code>MAP_NORESERVE</code>	No need to reserve space for the mapping.
<code>MAP_POPULATE</code>	Populate (prefault) page tables.
<code>MAP_NONBLOCK</code>	Do not block on I/O.

Creating an address interval

- On error `do_mmap()` returns a negative value
- On success
 - The kernel tries to merge the new interval with an adjacent one having same permissions
 - Otherwise, create a new VMA
 - Returns a pointer to the start of the mapped memory area
- `do_mmap()` is exported to user-space through `mmap2()`

```
void *mmap2(void *addr, size_t length, int prot,  
            int flags, int fd, off_t pgoffset);
```

Removing an address interval

- Removing an address interval is done through `do_munmap()`

```
/* linux/include/linux/mm.h */  
int do_munmap(struct mm_struct *, unsigned long, size_t);
```

- Exported to user-space through `munmap()`

```
int munmap(void *addr, size_t len);
```


Page tables

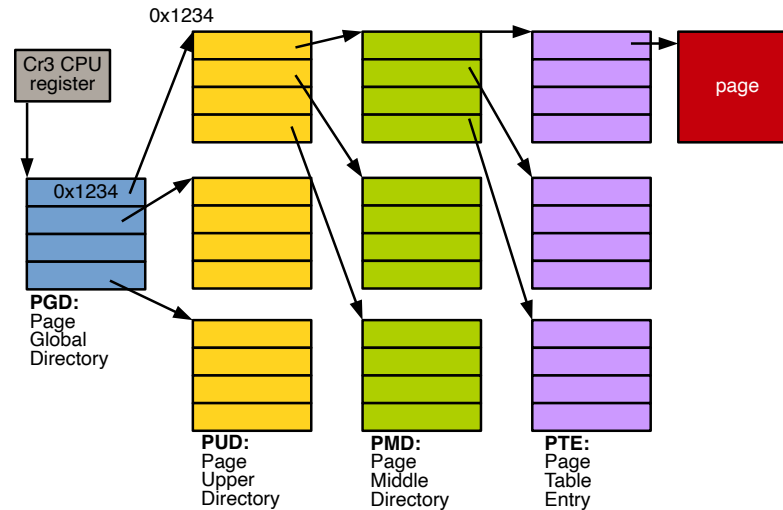
- Linux enables paging early in the boot process
 - All memory accesses made by the CPU are virtual and translated to physical addresses through the page tables
 - Linux set the page tables and the translation is made automatically by the hardware (MMU) according to the page tables content
- The address space is defined by VMAs and is sparsely populated
 - One address space per process → one page table per process
 - Lots of “empty” areas

Page tables

```

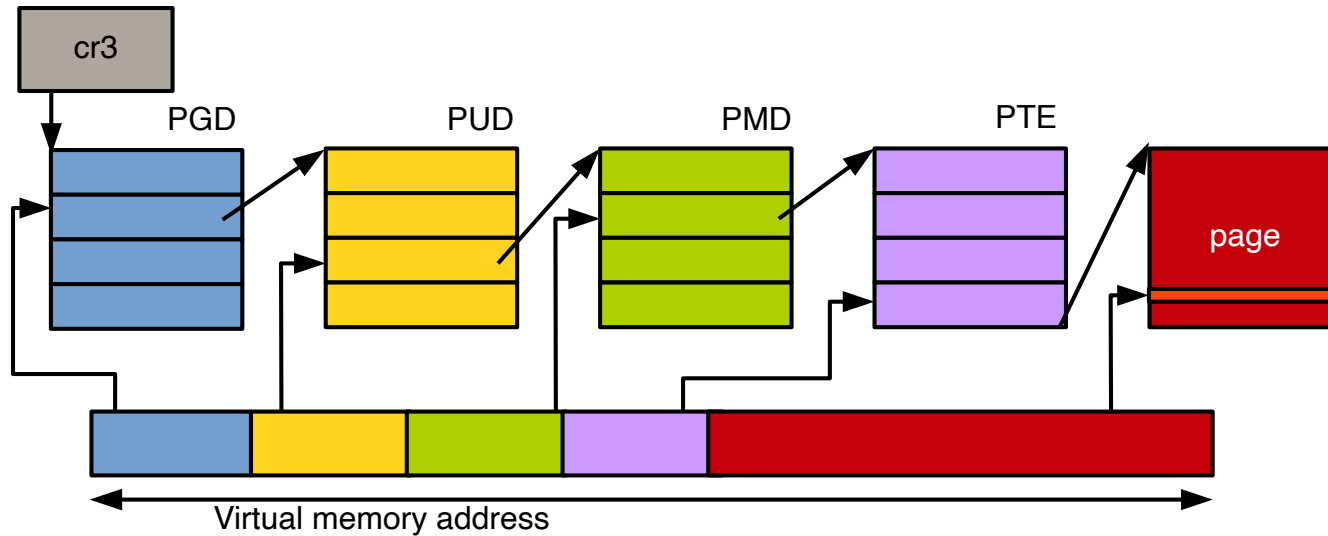
/* linux/include/linux/mm types.h */
struct mm_struct {
    struct vm_area_struct *mmap;           /* list of VMAs */
    struct rb_root          mm_rb;         /* rbtree of VMAs */
    pgd_t                   *pgd;          /* page global directory */
    /* ... */
};

```



Page tables

- Address translation is performed by the hardware (MMU)



Virtual address map in linux

```

0000000000000000 - 00007fffffffffff (=47 bits) user space, different per mm
hole caused by [47:63] sign extension
ffff800000000000 - ffff87fffffffffff (=43 bits) guard hole, reserved for hypervisor
ffff880000000000 - ffffc7fffffffffff (=64 TB)  **direct mapping of all phys. memory**
ffffc80000000000 - ffffc8fffffffffff (=40 bits) hole
ffffc90000000000 - ffffe8fffffffffff (=45 bits) vmalloc/ioremap space
ffffe90000000000 - ffffe9fffffffffff (=40 bits) hole
ffffea0000000000 - ffffeafffffffffffff (=40 bits) virtual memory map (1TB)
... unused hole ...
ffffec0000000000 - fffffbfffffffffff (=44 bits) kasan shadow memory (16TB)
... unused hole ...
                vaddr_end for KASLR
fffffe0000000000 - fffffe7fffffffffff (=39 bits) cpu_entry_area mapping
fffffe8000000000 - fffffefffffffffffff (=39 bits) LDT remap for PTI
ffffff0000000000 - ffffff7fffffffffff (=39 bits) %esp fixup stacks
... unused hole ...
fffffffef0000000 - ffffff9efffffffffffff (=64 GB)  EFI region mapping space
... unused hole ...
ffffffffff80000000 - fffffffffff9fffffffffff (=512 MB) kernel text mapping, from phys 0
ffffffffffa0000000 - fffffffffffefffffffffffff (1520 MB) module mapping space
[fixmap start] - fffffffffff5fffff kernel-internal fixmap range
ffffffffff600000 - fffffffffff600fff (=4 kB) legacy vsyscall ABI
fffffffffffe00000 - fffffffffffefffffffffffff (=2 MB) unused hole

```

Further readings

- [Introduction to Memory Management in Linux](#)
- [20 years of Linux virtual memory](#)
- [Linux Kernel Virtual Memory Map](#)
- [Kernel page-table isolation](#)
- [Addressing Meltdown and Spectre in the kernel](#)
- [Meltdown and Spectre](#)
- [Meltdown Attack Lab](#)

Further readings

- Supporting bigger and heterogeneous memory efficiently
 - [AutoNUMA](#), [Transparent Hugepage Support](#), [Five-level page tables](#)
 - [Heterogeneous memory management](#)
- Optimization for virtualization
 - [Kernel same-page merging \(KSM\)](#)
 - [MMU notifier](#)

Next class

- Virtual File System